



Non-Malleable Codes

STEFAN DZIEMBOWSKI, University of Warsaw

KRZYSZTOF PIETRZAK, Institute of Science and Technology (IST) Austria

DANIEL WICHS, Northeastern University

We introduce the notion of “non-malleable codes” which relaxes the notion of error correction and error detection. Informally, a code is non-malleable if the message contained in a modified codeword is either the original message, or a completely unrelated value. In contrast to error correction and error detection, non-malleability can be achieved for very rich classes of modifications.

We construct an efficient code that is non-malleable with respect to modifications that affect each bit of the codeword arbitrarily (i.e., leave it untouched, flip it, or set it to either 0 or 1), but independently of the value of the other bits of the codeword. Using the probabilistic method, we also show a very strong and general statement: there exists a non-malleable code for every “small enough” family \mathcal{F} of functions via which codewords can be modified. Although this probabilistic method argument does not directly yield efficient constructions, it gives us efficient non-malleable codes in the random-oracle model for very general classes of tampering functions—e.g., functions where every bit in the tampered codeword can depend arbitrarily on any 99% of the bits in the original codeword.

As an application of non-malleable codes, we show that they provide an elegant algorithmic solution to the task of protecting functionalities implemented in hardware (e.g., signature cards) against “tampering attacks.” In such attacks, the secret state of a physical system is tampered, in the hopes that future interaction with the modified system will reveal some secret information. This problem was previously studied in the work of Gennaro et al. in 2004 under the name “algorithmic tamper proof security” (ATP). We show that non-malleable codes can be used to achieve important improvements over the prior work. In particular, we show that any functionality can be made secure against a large class of tampering attacks, simply by encoding the secret state with a non-malleable code while it is stored in memory.

CCS Concepts: • **Mathematics of computing** → **Coding theory**; • **Security and privacy** → **Information-theoretic techniques**; *Tamper-proof and tamper-resistant designs*;

Additional Key Words and Phrases: Tamper resilience, provable security

ACM Reference format:

Stefan Dziembowski, Krzysztof Pietrzak, and Daniel Wichs. 2018. Non-Malleable Codes. *J. ACM* 65, 4, Article 20 (April 2018), 32 pages.

<https://doi.org/10.1145/3178432>

This work is supported by the National Natural Science Foundation under Grant No. CNS-1314722, and the European Research Council under Grants No. 207908-CNTM, No. 682815-TOCNeT, and No. 259668-PSPC.

Part of this work was done when Stefan Dziembowski was with *Sapienza* University of Rome, Italy.

Authors’ addresses: S. Dziembowski, University of Warsaw, Banacha 2, Warsaw, 02-097, Poland; email: S.Dziembowski@mimuw.edu.pl; K. Pietrzak, Institute of Science and Technology (IST) Austria, Am Campus 1, Klosterneuburg, Austria, 3400; email: pietrzak@ist.ac.at; D. Wichs, Northeastern University, 440 Huntington Avenue, Boston, Massachusetts, 02115; email: wichs@ccs.neu.edu.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2018 ACM 0004-5411/2018/04-ART20 \$15.00

<https://doi.org/10.1145/3178432>

1 INTRODUCTION

Consider the following three-step process, which we call the “tampering experiment”:

- (1) A source message s is encoded via a (possibly randomized) procedure Enc , yielding a *codeword* $c = \text{Enc}(s)$.
- (2) The codeword is modified under some *tampering function* $f \in \mathcal{F}$ to an *erroneous codeword* $\tilde{c} = f(c)$.
- (3) The erroneous codeword \tilde{c} is decoded using a procedure Dec , resulting in a *decoded message* $\tilde{s} = \text{Dec}(\tilde{c})$.

The tampering experiment can be used to model several interesting real-world settings, such as data transmitted over a noisy channel, or adversarial tampering of data stored in the memory of a physical device, which we will talk about more later. We would like to build special encoding/decoding procedures (Enc , Dec), which give us some meaningful guarantees about the results of the above tampering experiment, for large and interesting families \mathcal{F} of tampering functions. Let us explore several possibilities for the type of guarantees that we may hope for.

Error Correction. One very natural guarantee, called error correction, would be to require that for any tampering function $f \in \mathcal{F}$ and any source message s , the tampering experiment always produces the correct decoded message $\tilde{s} = s$. This notion of correction was first studied in the seminal work of Shannon in 1948 [67] with respect to random noise (e.g., each bit of the codeword c is flipped with some probability $p < 1/2$) and later by Hamming in 1950 [48], who introduced the notion of an *error-correcting code* with respect to worst-case errors. In particular, we can think of Hamming’s notion of an error-correcting code (with distance d) as guaranteeing the error-correction property for the family \mathcal{F} of functions f for which the *Hamming distance* between an erroneous codeword $\tilde{c} = f(c)$ and c is small (at most $d/2 - 1$). Since their introduction, error-correcting codes have received much attention and found countless applications in both theory and practice. There has also been some study of error correction for other families \mathcal{F} , such as ones where tampering functions preserve edit distance instead of Hamming distance (see the survey of [58]). However, it is also clear that error correction cannot be achieved for many natural (and simple) function families—such as even the single “zero function” f that maps all values x to the all-zero string.

Error Detection. A weaker guarantee, called error detection, requires that the tampering experiment always results in either the correct value $\tilde{s} = s$ or a special symbol $\tilde{s} = \perp$ indicating that tampering has been detected (where the latter can only happen when the codeword has been modified to $\tilde{c} \neq c$). This notion of error detection is a weaker guarantee than error correction, and achievable for larger families \mathcal{F} of tampering functions. For example, Hamming’s notion of an error-correcting code with distance d can detect up to $d - 1$ errors (but only correct $d/2 - 1$ errors).

The recent work of Cramer et al. [28] introduced the notion of “Algebraic Manipulation Detection” (AMD) codes. Such codes have an inherently *randomized* encoding procedure and guarantee that for *any a-priori chosen error vector* $\Delta \neq 0$, we get $\text{Dec}(\text{Enc}(s) + \Delta) = \perp$ with overwhelming probability over the randomness of the encoding.¹ In other words, these codes guarantee error detection with respect to the family \mathcal{F}_{err} consisting of *all constant-error functions* f_{Δ} given by $f_{\Delta}(x) \stackrel{\text{def}}{=} x + \Delta$. The definition of AMD codes is an interesting departure from the traditional study of error correction/detection, where restrictions are usually placed on the relationship between the codeword c and the erroneous codeword \tilde{c} (e.g., they are assumed to be close in some metric space). Under the tampering family \mathcal{F}_{err} , the erroneous codeword $\tilde{c} = f_{\Delta}(c) = c + \Delta$ can in principle take

¹We will often just use the addition operation $+$ over bit strings, where an n -bit string is interpreted as a value in \mathbb{F}_2^n .

on *any* possible value, and hence the final relationship between c and \tilde{c} can be arbitrary. However, the tampering functions $f_\Delta \in \mathcal{F}_{\text{err}}$ are *restricted in how they derive the erroneous codeword \tilde{c} from c* , since the error vector Δ is specified *a priori* and independently of the value of c (which contains some randomness from the encoding procedure). We take a similar view in this work, and study restrictions on the family \mathcal{F} of tampering functions rather than just the final relationship between a codeword and an erroneous codeword.

Unfortunately, even error detection cannot be achieved for many other natural (and simple) function families \mathcal{F} . For example, the family $\mathcal{F}_{\text{const}}$ of all *constant functions* $f_{c^*}(x) \stackrel{\text{def}}{=} c^*$, always contains some function that maps all values to a “valid” codeword c^* (for which $\text{Dec}(c^*) \neq \perp$), and hence is neither correctable nor detectable.

Non-Malleability. As the main focus of this work, we present yet another meaningful and previously unexplored notion, which we call “non-malleability.” A *non-malleable code* ensures that either the tampering experiment results in a correct decoded message $\tilde{s} = s$, or the decoded message \tilde{s} is *completely independent of and unrelated to* the source message s . In other words, our new notion of non-malleability for codes is similar, in spirit, to notions of non-malleability for cryptographic primitives (such as encryption,² commitments, and zero-knowledge proofs), introduced by the seminal work of Dolev et al. [34].

Compared to error correction or error detection, the “right” formalization of non-malleable codes is somewhat harder to define. Let Tamper_s^f be a random variable for the value of the decoded message \tilde{s} , which results when we run the tampering experiment with source-message s and tampering function f , over the randomness of the encoding procedure. Intuitively, we wish to say that the distribution of Tamper_s^f is independent of the encoded message s . Of course, we also want to allow for the case where the tampering experiment results in $\tilde{s} = s$ (e.g., if the tampering function is identity), which clearly depends on s . Thus, we require that for every tampering function $f \in \mathcal{F}$, there exists a distribution D_f which outputs either concrete values \tilde{s} or a special same symbol, and faithfully models the distribution of Tamper_s^f for all s in the following sense: for every source message s , the distributions of Tamper_s^f and D_f are statistically close when the same symbol is interpreted as s . That is, D_f correctly simulates the “outcome” of the tampering experiment with a function $f \in \mathcal{F}$ without knowing the source messages s , but it is allowed some ambiguity by outputting a same symbol to indicate that the decoded message should be the same as the source message, without specifying what the exact value is. The fact that D_f depends on *only* f and *not* on s , shows that the outcome of Tamper_s^f is independent of s , exempting equality.

Error correction w.r.t. to some class \mathcal{F} of tampering functions trivially implies non-malleability w.r.t. to \mathcal{F} : simply define D_f to always output the same symbol. Error detection only implies non-malleability if for every $f \in \mathcal{F}$, the probability $\Pr[\text{Dec}(f(\text{Enc}(s))) = \perp]$ (i.e., that a tampered codeword is invalid) is the same for all source message s . In that case, for each $f \in \mathcal{F}$ there is a distribution D_f over $\{\text{same}, \perp\}$ which satisfies the definition of non-malleability. For example, the above is the case for AMD codes where, for each $f_\Delta \in \mathcal{F}_{\text{err}}$, $\Pr[\text{Dec}(f_\Delta(\text{Enc}(s))) = \perp] = \Pr[\text{Dec}(\text{Enc}(s) + \Delta) = \perp]$ is either 0 if $\Delta = 0$, or negligibly close to 1 if $\Delta \neq 0$, independently of s .

1.1 Main Results for Non-Malleability

It is easy to see that non-malleability cannot be achieved w.r.t. to all functions; for example, consider a class \mathcal{F} consisting of only one function f which flips all bits in the encoded message (i.e.,

²For example, a non-malleable encryption scheme ensures that, by modifying the ciphertext in any way, the decrypted value is unrelated to that contained in the original ciphertext.

$f(c) = \text{Enc}(\text{Dec}(c) \oplus 1^k)$). Moreover, if one requires that Enc and Dec are computable in polynomial time, then non-malleability cannot be achieved w.r.t. to all *poly-time computable* functions (as in this case f defined above is also poly-time computable). Therefore, one needs to consider restricted classes of tampering functions. In this article, we show that non-malleability can be achieved for many rich families \mathcal{F} of tampering functions (for which error correction and error detection are impossible). While the fact that non-malleability cannot be achieved with respect of all functions may be slightly disappointing, we would like to stress that our restrictions correspond to many realistic tampering scenarios (like the bitwise tampering or the split-state tampering). Moreover, as shown in the subsequent work, non-malleable codes for several restricted families can be useful for constructing other cryptographic primitives, like the encryption secure against the chosen-ciphertext attacks [25] and the string commitment schemes [5].

Bit-wise Independent Tampering. As one very concrete example, we study non-malleability with respect to the family of functions f which specify, for each bit of the codeword c , whether to keep it as is, flip it, set it to 0, or set it to 1. That is, each bit of the codeword is modified *arbitrarily* but independently of the value of the other bits of the codeword. We call this the “bit-wise independent tampering” family \mathcal{F}_{BIT} . Note that this family contains constant functions $\mathcal{F}_{\text{const}}$ and constant-error functions \mathcal{F}_{err} as subsets. Therefore, as we have mentioned, error correction and error detection cannot be achieved w.r.t. this family. Nevertheless, we show an efficient non-malleable code for this powerful family. Our construction, described in Section 4, is based on AMD codes (described above) and a type of linear secret-sharing scheme (over bits) with some additional properties.

All Families of Bounded Size. Even non-malleability cannot be achieved with respect to the family of *all* tampering functions. No matter what encoding/decoding procedure one chooses, there is always some function f , which, on input c , computes $s = \text{Dec}(c)$, sets s' to be the same as s but with, for example, the last bit flipped, and outputs $\tilde{c} = \text{Enc}(s')$. However, we notice that this “bad” function f depends on the encoding/decoding procedure. Therefore, even though *for any* coding scheme *there exists* some tampering function f that breaks non-malleability, we show that *for any* “small enough” tampering family \mathcal{F} , *there exists* some coding scheme which is non-malleable w.r.t. \mathcal{F} . Note that, when considering codewords of size n , the family \mathcal{F}_{all} of all tampering functions on n -bit codewords has size given by $\log(\log(|\mathcal{F}_{\text{all}}|)) = n + \log(n)$. Our result shows that for any *slightly smaller* family \mathcal{F} whose size is given by $\log(\log(|\mathcal{F}|)) < n$, there exist non-malleable codes w.r.t. \mathcal{F} and, in particular, a random code is likely to be non-malleable with overwhelming probability. This is in contrast to error correction and error detection, for which there exist some small families \mathcal{F} (of size $\log(\log(|\mathcal{F}|)) = \log(n)$) for which *no* code can be error correcting or error detecting (e.g., the family $\mathcal{F}_{\text{const}}$).

Unfortunately, our existential result is derived via a probabilistic method argument, and therefore does not directly lead to efficient constructions. However, as a corollary of the probabilistic method argument, we can show that efficient codes in the “random oracle model” are non-malleable w.r.t. large function families \mathcal{F} . For example, we get such non-malleable code in the split-state model, that is, a code that is secure with respect to all functions which the codeword is divided into two parts with which the adversary can tamper independently. More formally, the *split-state* family consists of functions $f : \{0, 1\}^n \rightarrow \{0, 1\}^n$ which can be written as f_1, f_2 for $f_1, f_2 : \{0, 1\}^{n/2} \rightarrow \{0, 1\}^{n/2}$, such that, for any value $c = (c_1, c_2)$, we have $f(c_1, c_2) = (f_1(c_1), f_2(c_2))$. Note that this is reminiscent of the independent leakage model used often in the leakage-resilient cryptography [30, 37, 39, 47]. More generally, we get such codes for the family of functions f , where the i th bit in $f(c)$ is a function of an arbitrary fraction (say, 99%) of the bits of c .

1.2 Application to Tamper-Resilient Cryptography

Traditionally, cryptographic security notions assume that an adversary only has “black-box” access to an attacked system. That is, she can only observe the input/output behavior to the system, but gets no further information. Such models often fail to capture physical reality, where an adversary can attack an actual physical implementation of the system. Attacks which go beyond black-box access are called “implementation attacks,” and we distinguish between two classes of such attacks, namely, “leakage” and “tampering” attacks.³ Leakage attacks refer to all attacks where the adversary can learn some information about the secret internal state of a system by measuring some properties of its physical implementation (e.g., timing, radiation, heat, power consumption). Tampering attacks, on the other hand, refer to attacks where the adversary actively modifies the computation (e.g., by cutting wires in the circuit or heating it to introduce random errors in memory), and then learns some information from the input/output behavior of the modified computation. As an example of how such attacks are used, the work of Boneh et al. [12] shows that if random faults are introduced into the computation of a single RSA signature, then the resulting output can be used to factor the modulus. Of course, a physical attacker can combine both leakage and tampering attacks.

To get any non-trivial security for an implementation, some degree of security against both leakage and tampering attacks is necessary: no security can be achieved if the device just leaks the entire secret state, and no security can be achieved if the adversary can tamper the computation to the extent that she can simply replace the original computation with a computation that outputs the entire secret state. The cryptographic community has recently seen a flurry of work on cryptographic systems with leakage-resilience properties [7, 8, 13, 31–33, 35–37, 39, 41, 44, 49, 51, 53, 57, 59, 63]. On the other hand, we are aware of only two works which formally study security against tampering attacks: the work of Genaro et al. [45] and Ishai et al. [50].⁴ Our model of tamper resilience is similar to that of [45] and we first describe this model and our results. We then compare this to the results (and model) of [45, 50].

Model of Tamper-Resilient Security. Following Gennaro et al. [45], we consider the problem of constructing a secure device by combining secret “leakage-proof” but *not* “tamper-proof” memory with a public “tamper-proof” and “leakage-proof” circuit. Thus, only the parts of the device that are fixed and publicly known cannot be tampered with. The advantage of this approach is that one only has to manufacture some fixed tamper-proof circuitry, but the memory—which must be read/written constantly and thus is harder to protect—need not be tamper proof. On the other hand, as discussed in [44, 57], the requirement that all components of the device are completely “leakage proof” is a very strong one and very hard, if not impossible, to satisfy in practice. Luckily, we can use the above-mentioned recent results on leakage-resilient cryptography to weaken this requirement, and allow some bounded amount of information to leak during computation. Therefore, we are left with the task of securing the tamper-prone memory against tampering attacks.

Our Results. We consider an arbitrary system $\langle G, s \rangle$ consisting of a *public functionality* G assumed to be implemented on a “tamper-proof” and “leakage-proof” circuit, and a *secret state* s stored in “leakage-proof” but “tamper-prone” memory. The system may be stateful and interactive: a user can invoke it periodically with inputs x by submitting an $\text{Execute}(x)$ command, at which point the system computes $(y, s') \leftarrow G(x, s)$, updates its state to $s := s'$, and outputs y . We call this type of

³Leakage attacks are often called “passive attacks,” whereas tampering (or tampering combined with leakage) attacks are called “active attacks.”

⁴Ad hoc solutions for some particular systems were proposed by [46, 56]. Some of these were subsequently broken [19, 26], highlighting the need for formal study.

interaction with the system *black-box access*. We also assume that an adversary can, in addition, interact with the system by issuing $\text{Tamper}(f)$ commands for some tampering functions $f \in \mathcal{F}$, at which point the system state is modified to $s := f(s)$. We call this type of interaction (via both Execute and Tamper commands) *tampering access* to the system. Our goal is to take any system $\langle G, s \rangle$ and compile it into a new system $\langle G', s' \rangle$, which acts exactly the same as the original in any black-box interaction (i.e., preserves functionality), but is also *tamper resilient*. We formalize the latter by requiring that, for any adversary \mathcal{A} with tampering access to the compiled system $\langle G', s' \rangle$, there is a simulator \mathcal{S} which only has black-box access to the original system $\langle G, s \rangle$ and “learns” the same information as \mathcal{A} . In other words, tampering access in the compiled system does not give \mathcal{A} any advantage over black-box access to the original.

As our main result for tamper-proof security, we show that any coding scheme (Enc, Dec) , which is non-malleable w.r.t. some family \mathcal{F} , can be used to compile *any* system $\langle G, s \rangle$ into a system $\langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle$ which is tamper resilient w.r.t. \mathcal{F} . The compiled system $\langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle$ stores its state in encoded form. As the first step in any computation, the compiled system decodes its state and then runs the original computation with the decoded state and the input. Lastly, at the end of the computation, the compiled system re-encodes its updated state and updates the memory contents with the new state. On a high level, the compiled system is tamper resilient since tampering with the encoded state c via a $f \in \mathcal{F}$ produces a predictable result, which is independent of the underlying state s . More precisely, the tampering interaction can be simulated by sampling from the distribution D_f used to define non-malleability: if the outcome is the same symbol, tampering access to the compiled system can be simulated by black-box access to the original system, and, if the outcome is some specific value \tilde{s} , tampering access can be simulated by running $\langle G, \tilde{s} \rangle$ independently of the original system.

Comparison to “Algorithmic Tamper-Proof Security”. As we mentioned, our model of tamper-resilient security, where we assume a “tamper-proof”/“leakage-proof” public circuit and a “leakage-proof” but *not* “tamper-proof” memory was proposed and studied by Gennaro et al. [45]. The main solution of that work consists of using a (strong) *public-key signature scheme* and storing the state s along with a signature $\sigma = \text{Sign}_{sk}(s)$ on memory. In our context, we can think of this solution as using a *keyed encoding scheme* $(\text{Enc}_{sk}, \text{Dec}_{pk})$ where $\text{Enc}_{sk}(s)$ is now a *secret encoding algorithm* which outputs s along with a signature $\sigma = \text{Sign}_{sk}(s)$ under the secret-key sk . The decoding function $\text{Dec}_{pk}(s, \sigma)$ verifies the signature σ using pk , outputs \perp if the signature is invalid, and s otherwise. We compare the above solution with our solution using non-malleable codes.

Limited Functionality. The security notion achieved by the signature-based coding scheme does not satisfy our notion of non-malleability, even for simple/natural function families like bitwise independent tampering \mathcal{F}_{BIT} . The problem is that a tampering attacker can learn some limited information about the state s by, for example, setting the first bit of the encoded state (s, σ) to 0. Then the decoded message is s *if and only if* the first bit of s already was 0, and \perp otherwise (as w.h.p. σ will not be a valid signature). Therefore, the adversary can learn some bits of s via tampering attacks.⁵ This issue was already recognized in the work of [45], but it was noted that tamper-resilient security is still achieved for some functionalities such as *signatures and encryption schemes* where leaking logarithmically many bits about the secret key can be tolerated without breaking the security. Unfortunately, even in these special cases, some problems come up: for example, security breaks down if an adversary has tampering access to *many* signature/encryption devices that use

⁵In fact, using this type of an attack, one can learn up to logarithmically many bits of s with polynomial probability. A “self-destruct” feature, which we will discuss below, will prevent the adversary from launching this attack many times on different bits of the state, ultimately learning the entire secret.

the same secret key, and encryption security breaks down if the adversary can adaptively perform a tampering attack based on the challenge ciphertext. More generally, the signature-based solution does not provide a general method for making any functionality G tamper resilient. In contrast, our solution using non-malleable codes gives us strong guarantees *for all functionalities* G , ensuring that a tampering attack can always be simulated.

Secret-keyed Encoding. Another big difference between our solution and that of [45] is that our encoding/decoding procedures are public and un-keyed, while that of [45] requires a secret-keyed encoding procedure (and a public-keyed decoding). This has several implications. Firstly, the use of a secret-keyed encoding means that state of G has to be decided by an outside party (that knows the secret signing key sk) and cannot be done by the system itself. In particular, this solution will *not* work for stateful functionalities G which need to update their own state. Secondly, the use of keys brings up some key-management problems where the “software” designer (who decides on the state s) needs to coordinate with the hardware designer (who knows the signing key sk). In contrast, our solution has public encoding/decoding and therefore is applicable to stateful functionalities G without requiring key management.

Assumptions. The signature-based solution requires computational assumptions (one-way functions), while our main constructions of non-malleable codes are information theoretic.

Tampering Family \mathcal{F} . The main benefit of the signature-based solution is that it allows for the most general family \mathcal{F} of *all efficient tampering functions*. In contrast, all of our solutions (necessarily) only consider smaller classes of tampering functions, namely, those for which non-malleable codes exist. In this aspect, our result is weaker but still well motivated since tampering attacks are unlikely to be “too complicated.” To the best of our knowledge, bitwise independent tampering (considered in Section 4) already covers the majority of real-world tampering attacks that have been demonstrated in practice.

Self-Destruct/Self-Update. The ATP model of [45] assumes a “self-destruct” feature, which means that whenever the memory content decodes to \perp (i.e., is not of the form (s, σ) where σ is a valid signature of s), the device is “destroyed.” This can be done by overwriting the memory content with some dummy value, say the all-zero string, or setting some (tamper-proof) flag that prevents future use of the device.

Our construction of a tamper-proof system via non-malleable codes has a similar self-destruct feature implicitly built into it: the system freshly re-encodes its state after each invocation, where, if the decoding gives \perp , the re-encoded value is set to a dummy all-zero string. For functionalities $G(., .)$ which do not update their state, this requirement imposes some overhead as memory must be written to.

It would be desirable to have a transformation achieving tamper-proof security w.r.t. to some class \mathcal{F} from codes which are non-malleable w.r.t. \mathcal{F} (or some similar notion), where also the transformed system is stateless. We leave this problem for future work. Here, we would only like to mention that our construction of non-malleable codes for bitwise independent tampering (\mathcal{F}_{BIT}) has some particular properties⁶ which allow us to prove that our transformation (applied to stateless functionalities) remains tamper proof even if we simply omit updating the state. (We then need an explicit self-destruct as described above.)

⁶Informally, given some history f_1, \dots, f_{i-1} of tampering functions that were applied to some unknown codeword, we can compute (assuming the codeword is still valid) the probabilities that applying a tampering function f_i will (1) produce an invalid codeword, (2) leave the encoded message unchanged, or (3) produce an encoding of a different message. In the last case we can also output a codeword having the right distribution.

Comparison to “Private Circuits II”. In [50] the authors build a general circuit compiler where an adversary, who can tamper with the wires of a compiled circuit, does not learn any more information than an adversary that only has black-box access to the original circuit. The transformation uses techniques from multi-party computation and is quite delicate. The model considered in that work is stronger than ours as it allows tampering of the “circuit” and “memory” (essentially, memory corresponds to a subset of the wires associated with the state), whereas we only consider tampering of “memory” but assume the circuit is tamper proof. However, in this very general setting, the work of [50] only considers a very limited tampering-function family that modifies a bounded subset of the wires between each invocation (to tolerate tampering of t wires per invocation, the compiler will blow up the circuit size by a factor of at least t). When restricted to memory, this would be analogous to standard error-correcting codes which provide security against a bounded number of individual bit modifications. In contrast, we achieve security for much richer families \mathcal{F} of tampering functions. It is an interesting open problem to combine the two approaches and achieve tamper-proof security for circuits as [50], but for larger families \mathcal{F} , and, ideally for all families for which there exist non-malleable codes.

1.3 Subsequent Work

After the conference version of this article was published [40], the notion of non-malleable codes attracted significant attention from the research community, and several papers on this topic have been published. In this section, we briefly summarize this subsequent work.

Constructions in the Split-state and Blockwise Tampering Models. A significant effort has been devoted to constructing non-malleable codes in the split-state model (cf. Section 1.1) without relying on the random oracle assumption. The first result in this sequence of papers was [54] where the authors constructed such a code in the common reference string model using the computational assumptions. An information-theoretically secure construction (in the plain model) for messages of length 1 was given in [38]. The first construction of such codes for messages of arbitrary length was shown in [3]. The codeword c in this construction is of length $O(|s|^7 \log |s|)$, where $|s|$ is the length of the encoded message (this was later improved by [1] to $O(|s|^6)$). The first construction of information-theoretically secure non-malleable codes in the split-state model with linear rate (i.e., such that $|c| = O(|s|)$) was given in [20], at the price of increasing the number of memory parts to 10. A linear-rate construction on two memory parts was recently shown in [2].

Also very recently, Chandran et al. [18] proposed an interesting generalization of the split-state model that they call the *blockwise tampering*. Roughly speaking, in this model the codeword is divided into a number of blocks (c_1, \dots, c_n) and the tampering function for each i th block can “look at” the previous blocks c_1, \dots, c_{i-1} (moreover, the “order of blocks” can be chosen by the adversary; we refer the reader to [18] for more on this model). The authors show that no information-theoretically secure non-malleable codes exist in this model, and they provide a computationally-secure construction.

Affine Mauling Functions. In another interesting model, the codeword is interpreted as a vector from some linear space \mathbb{F}^n (over a finite field \mathbb{F}), and the tampering function is an affine transformation over this space. [17] shows an encoding that is non-malleable against a subclass of affine tampering functions where $\mathbb{F} = \mathbb{Z}_2$. This is achieved using the results from the area of secure network coding. Also, the split-state encoding of [38] is secure against the affine mauling functions (as shown there, in Section 6).

Bitwise Tampering, Permutations, and Perturbations. An alternative (to ours) construction of non-malleable codes secure against bitwise tampering was given in [16]. This is achieved by

exploiting a connection between this notion and the wiretap channel of [62]). A construction with improved parameters was given in [23]. A natural generalization of this model is to allow the tampering functions to permute the bits of the codeword. Codes that are non-malleable with respect to such functions were constructed in [5, 6].

Improved Probabilistic Method Results. In this work, we employ the probabilistic method argument to show the existence of inefficient non-malleable codes for any tampering family \mathcal{F} of sufficiently small size $\log \log |\mathcal{F}| < n$. Subsequent work improved this argument in two important directions. Firstly, the work of [22] explores the rate (defined as the ratio of message size to codeword size) of non-malleable codes, showing that for function families satisfying $\log \log |\mathcal{F}| = \alpha n$ for some constant $\alpha < 1$, the rate can be as high as $1 - \alpha$, which is optimal in general. (In contrast, the result in our work only yields a sub-optimal rate of $(1 - \alpha)/3$.) Secondly, the works of [22, 43, 52] show that one can rely on a more careful probabilistic method argument with bounded independence to show the existence of *efficient* non-malleable codes for tampering families \mathcal{F} of singly exponential size $\log |\mathcal{F}| = s(n)$ for some polynomial $s(n)$. However, instead of specifying a fixed code which is guaranteed to be non-malleable for such family \mathcal{F} , these results give a family of efficient codes such that a random member of the family is non-malleable for \mathcal{F} with overwhelming probability.

Extensions of the Non-malleability Notion. The notion of non-malleable codes was extended in order to capture also other types of physical attacks. Non-malleable codes that are secure against the information leakage were constructed in [54] (in the computationally secure variant) and in [4, 29, 38]. Codes that are non-malleable to *continual* tampering were constructed in [42, 52].

Locally decodable and updatable non-malleable codes were constructed in [29].

Connection to Other Notions and Applications of the Non-malleable Codes. Some authors have also studied connections between the non-malleable codes and the other notions in cryptography. In [23], the authors explore a connection between the non-malleable codes in the split-state model and the seedless non-malleable extractors (which is a new notion that they introduce). In [52], the authors propose a new notion of *tamper-detection* codes and show how to construct the non-malleable codes out of the tamper-detection ones.

The non-malleable codes turned out to be a useful tool for solving the *domain extension problem*, that is, constructing a given cryptographic primitive for longer messages from the same primitive that works for short messages. This was done for the non-malleable commitments in [5], and for various flavors of public-key encryption in [24, 25]. Non malleable codes were used to construct commitment schemes also in [18].

[29] use the non-malleable codes for securing RAM computation against memory tampering and leakage attacks.

2 PRELIMINARIES AND NOTATION

For a randomized function g , we write $g(x; r)$ to denote the value of g on input x using randomness r . Sometimes we do not want to make the randomness explicit and only write $g(x)$ to denote the random variable $g(x; R)$ for a uniformly random R . For a random variable X , we use the notation $x \leftarrow X$ to denote that x is sampled randomly according to X . For a set S , we use the notation $s \leftarrow S$ to denote that s is sampled uniformly at random from S .

The *Hamming weight* of a binary string $x \in \{0, 1\}^n$, denoted $w_H(x)$, is the number of 1's in x . The *Hamming distance* of two strings x, x' , denoted $d_H(x, x') \stackrel{\text{def}}{=} w_H(x - x')$ is the number of positions in which they differ.

The *Statistical Distance* of two random variables X_0, X_1 with support \mathcal{X} is

$$\text{SD}(X_0, X_1) \stackrel{\text{def}}{=} \frac{1}{2} \cdot \sum_{x \in \mathcal{X}} |\Pr[X_0 = x] - \Pr[X_1 = x]|.$$

If the statistical distance between X_0 and X_1 is negligible (as a function of some well-defined security parameter), we say that X_0 and X_1 are statistically indistinguishable and write $X_0 \approx X_1$. We write $X_0 \equiv X_1$ if they are identically distributed.

We will frequently describe random variables in forms of “experiments” (see, e.g., Definition 3.2) which will have a form of a sequence $\{inst_1, \dots, inst_n\}$ of *instructions*. Every instruction $inst_i$ will be either of a form $x \leftarrow \mathcal{X}$ (meaning: variable x is sampled uniformly from set \mathcal{X}), or of a form $x \leftarrow \text{Alg}(x_1, \dots, x_k)$, where x_1, \dots, x_k are variables that were defined in earlier instructions $inst_1, \dots, inst_{i-1}$, and Alg is a randomized algorithm. The meaning of this is, algorithm Alg is run on x_1, \dots, x_k and some uniformly sampled internal randomness, and then the result of this computation is assigned to variable x . The final instruction $inst_n$ has a form “Output w ,” where w is some expression that involves the variables that were defined earlier in the experiment. The random variable defined by the experiment $\{inst_1, \dots, inst_n\}$ takes value of w (where the randomness is taken over the random choice done in uniform sampling and the choice of the internal randomness of the algorithms).

A sequence C_1, \dots, C_n of random variables is *t-wise independent* if every subsequence of size t of these variables (i.e., C_{i_1}, \dots, C_{i_t} for some i_1, \dots, i_t) is independent.

In this article, we only consider binary logarithms, and hence \log will denote the logarithm to the base 2.

In Section 5.2, we use the *random-oracle model* [10]. Formally, an *oracle* \mathcal{O} is a machine that is parametrized by a function $h : \{0, 1\}^* \rightarrow \{0, 1\}^t$, for some parameter t called the *output length* of \mathcal{O} . We say that \mathcal{O} is a *random oracle* if h is chosen uniformly at random from the set of all functions from $\{0, 1\}^*$ to $\{0, 1\}^t$. The oracle \mathcal{O} can be queried on inputs $x \in \{0, 1\}^*$, and it replies to each such query with $h(x)$. A machine M is called an *oracle machine* if it is an interactive Turing machine that interacts with an oracle. For simplicity, we consider only oracle machines that always produce some output. The function computed by M after interacting with an oracle \mathcal{O} will be denoted as $M^{\mathcal{O}}$.

3 NON-MALLEABLE CODES

Definition 3.1 (Coding Scheme). A coding scheme consists of two functions: a *randomized* encoding function $\text{Enc} : \{0, 1\}^k \rightarrow \{0, 1\}^n$, and a *deterministic* decoding function $\text{Dec} : \{0, 1\}^n \rightarrow \{0, 1\}^k \cup \{\perp\}$ such that, for each $s \in \{0, 1\}^k$, $\Pr[\text{Dec}(\text{Enc}(s)) = s] = 1$ (over the randomness of the encoding algorithm).

We can also talk about coding-scheme ensembles, parametrized by a message length $k \in \mathbb{N}$ and security parameter $\lambda \in \mathbb{N}$, so that, for each choice of k, λ , the ensemble contains an efficient code with $n = n(k, \lambda)$. We will often assume such ensembles implicitly, but will try not to complicate the exposition with additional notation.

We now define non-malleability with respect to some family \mathcal{F} of tampering functions. We want to say that a code is non-malleable, if the result of tampering with a codeword is independent of the encoded message. To do so, we require that for each $f \in \mathcal{F}$ there is a universal distribution D_f which, for all s , indicates what the likely outcomes are when applying a tampering function f to a (random) encoding of s . The distribution D_f only gets (black-box) access to $f(\cdot)$, but does not get s . We allow D_f to output a special same symbol, to indicate that tampering via f does not change the initial encoded message (without needing to commit to what that message is).

Definition 3.2 (Non-Malleability). Let \mathcal{F} be some family of tampering functions. For each $f \in \mathcal{F}$ and $s \in \{0, 1\}^k$, define the tampering experiment

$$\text{Tamper}_s^f \stackrel{\text{def}}{=} \left\{ \begin{array}{l} c \leftarrow \text{Enc}(s), \tilde{c} \leftarrow f(c), \tilde{s} = \text{Dec}(\tilde{c}) \\ \text{Output: } \tilde{s}. \end{array} \right\}$$

which is a random variable over the randomness of the encoding function Enc . We say that a coding scheme (Enc, Dec) is *non-malleable with respect to \mathcal{F}* if for each $f \in \mathcal{F}$, there exists a distribution D_f over $\{0, 1\}^k \cup \{\perp, \text{same}\}$, such that, for all $s \in \{0, 1\}^k$, we have

$$\text{Tamper}_s^f \approx \left\{ \begin{array}{l} \tilde{s} \leftarrow D_f \\ \text{Output } s \text{ if } \tilde{s} = \text{same}, \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\} \quad (1)$$

and D_f is efficiently samplable given oracle access to $f(\cdot)$. Here “ \approx ” can refer to statistical or computational indistinguishability. In the case of statistical indistinguishability, we say the scheme has error ε , if the statistical distance above is at most ε .

We also define a slightly stronger notion, called *strong non-malleability*, which essentially says that, whenever the codeword is actually modified to $\tilde{c} \neq c$, the decoded message \tilde{s} is independent of s . This is in contrast to the plain notion of non-malleability where some modifications of the codeword c could preserve the value of the contained message $\tilde{s} = s$. The stronger notion of non-malleability is somewhat simple to define.

Definition 3.3 (Strong Non-Malleability). Let \mathcal{F} be a family of functions. We say that a coding scheme (Enc, Dec) is *strongly non-malleable with respect to \mathcal{F}* if for any $s_0, s_1 \in \{0, 1\}^k$ and any $f \in \mathcal{F}$, we have

$$\begin{aligned} \text{StrongNM}_{s_0}^f &\approx \text{StrongNM}_{s_1}^f, \\ \text{where } \text{StrongNM}_s^f &\stackrel{\text{def}}{=} \left\{ \begin{array}{l} c \leftarrow \text{Enc}(s), \tilde{c} \leftarrow f(c), \tilde{s} = \text{Dec}(\tilde{c}) \\ \text{Output } \text{same} \text{ if } \tilde{c} = c, \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\} \end{aligned} \quad (2)$$

Here “ \approx ” can refer to statistical, or computational indistinguishability. In the case of statistical indistinguishability, we say the scheme has error ε , if the statistical distance is at most ε .

As implied by our naming, strong non-malleability is a strictly stronger property than (regular) non-malleability.

THEOREM 3.1. *If a coding scheme (Enc, Dec) is strongly non-malleable with respect to some family \mathcal{F} , then it is non-malleable with respect to \mathcal{F} with the same error.*

PROOF. Set $D_f = \text{StrongNM}_{s_0}^f$ for some constant $s_0 \in \{0, 1\}^k$. Then, for each $s \in \{0, 1\}^k$, $D_f \approx \text{StrongNM}_s^f$ and so

$$\begin{aligned} &\left\{ \begin{array}{l} \tilde{s} \leftarrow D_f \\ \text{Output } s \text{ if } \tilde{s} = \text{same}, \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\} \\ &\approx \left\{ \begin{array}{l} \tilde{s} \leftarrow \text{StrongNM}_s^f \\ \text{Output } s \text{ if } \tilde{s} = \text{same}, \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\} \\ &\equiv \left\{ \begin{array}{l} c \leftarrow \text{Enc}(s), \tilde{c} \leftarrow f(c), \tilde{s} = \text{Dec}(\tilde{c}) \\ \text{Output: } \tilde{s}. \end{array} \right\}. \quad \square \end{aligned}$$

We also note that our notion of *strong* non-malleability bears a very close resemblance to the standard notion of non-malleable encryption as defined in [34]. On the other hand, our default notion of non-malleability (where tampering may leave the message unchanged) is closer, in analogy,

to a relaxed notion of cryptographic non-malleability called “replayable” non-malleable security, defined in [15]. Indeed, in that work it was already noted that for many cryptographic applications this relaxed notion of non-malleable encryption suffices. When it comes to non-malleable codes, we chose to make the relaxed notion the default one since

- Having a code which ensures that tampering *preserves* the message entirely seems like a desirable feature rather than a bug when designing coding schemes. In particular, we view non-malleability as a weakening of error correction. This view is only valid with the relaxed notion of non-malleability, since error correction does *not* imply *strong non-malleability*.
- Our default relaxed notion of non-malleability is sufficient for our application to tamper resilience.

The main disadvantage of our default notion of non-malleability (compared to strong non-malleability) seems to be a more complicated definition, which is harder to state and understand. Inspired by [15], we also give an alternative definition of our default notion of non-malleability in Appendix A, and prove the equivalence of the two definitions.

4 BITWISE INDEPENDENT TAMPERING

With \mathcal{F}_{BIT} , we denote the family which contains all tampering functions that tamper every bit independently. Formally, this family contains all functions $f : \{0, 1\}^n \rightarrow \{0, 1\}^n$ that are defined by n functions $f_i : \{0, 1\} \rightarrow \{0, 1\}$ (for $i = 1, \dots, n$) as $f(c_1, \dots, c_n) = (f_1(c_1), \dots, f_n(c_n))$. Note that there are only four possible choices for each f_i (i.e., how to modify a particular bit) and we name these “set to 0,” “set to 1,” “flip,” and “keep” where the meanings should be intuitive. We call the above family the *bitwise independent tampering family*.

As discussed in the Introduction, standard notions of “error correction” and “error detection” trivially cannot be achieved against this class. Note that the family \mathcal{F}_{err} of constant-error functions which was studied in the context of AMD codes is a subset of \mathcal{F}_{BIT} . More concretely, \mathcal{F}_{err} is defined as a set of functions in \mathcal{F}_{BIT} , except that f_i can only be “keep” or “flip,” but not “set to 0” or “set to 1.” Bitwise independent tampering is interesting in the context of tamper-proof security, as most practical attacks tamper with each bit in the computing circuit and/or the memory individually. In Section 6, we will show how non-malleable codes can be used to protect any functionality against tampering of the memory.

4.1 Construction

Before proceeding with the construction, we first define two primitives which are interesting in their own right: a type of secret-sharing scheme over bits that also has a large distance, and AMD codes which we mentioned previously in the Introduction.

Definition 4.1 (LECSS Schemes). Let $(E_{\text{SS}}, D_{\text{SS}})$ be a coding scheme with source messages $s \in \{0, 1\}^k$ and codewords $c \in \{0, 1\}^n$. We say that the scheme is a (d, t) -linear error-correcting secret-sharing (LECSS) scheme if the following properties hold:

- *Linearity:* For all $c \in \{0, 1\}^n$, such that $D_{\text{SS}}(c) \neq \perp$, all $\Delta \in \{0, 1\}^n$, we have

$$D_{\text{SS}}(c + \Delta) = \begin{cases} \perp & \text{if } D_{\text{SS}}(\Delta) = \perp \\ D_{\text{SS}}(c) + D_{\text{SS}}(\Delta) & \text{otherwise.} \end{cases}$$

- *Distance d :* For all non-zero $\tilde{c} \in \{0, 1\}^n$ with hamming weight $w_H(\tilde{c}) < d$, we have $D_{\text{SS}}(\tilde{c}) = \perp$.
- *Secrecy t :* For any fixed $s \in \{0, 1\}^k$, we define the random variables $C = (C_1, \dots, C_n) = E_{\text{SS}}(s)$, where C_i denotes the bit of C in position i (and where the randomness comes from

the encoding procedure). Then the random variables $\{C_i\}_{1 \leq i \leq n}$ are individually uniform over $\{0, 1\}$ and t -wise independent.

Definition 4.2 (AMD Codes [28]). Let $(E_{\text{AMD}}, D_{\text{AMD}})$ be a coding scheme with $E_{\text{AMD}} : \{0, 1\}^k \rightarrow \{0, 1\}^n$. We say that $(E_{\text{AMD}}, D_{\text{AMD}})$ is a ρ -secure AMD code if for all $s \in \{0, 1\}^k$ and all non-zero $\Delta \in \{0, 1\}^n$ we have $\Pr[D_{\text{AMD}}(E_{\text{AMD}}(s) + \Delta) \neq \perp] \leq \rho$, where the probability is over the randomness of the encoding.

The definition of AMD codes given here is slightly stronger than the one from Cramer et al. [28] which required $\Pr[D_{\text{AMD}}(E_{\text{AMD}}(s) + \Delta) \notin \{s, \perp\}] \leq \rho$. In particular, the weaker definition allows for a non-zero Δ to leave the message unchanged. However, the construction of AMD codes from Cramer et al. [28] satisfies our stronger definition; indeed the proof given in that work already shows this without requiring any modification.⁷

Our main theorem of this section shows that, by composing an AMD code with a LECSS scheme, we get a non-malleable code for the bitwise independent tampering family \mathcal{F}_{BIT} .

THEOREM 4.1. *Let $(E_{\text{SS}}, D_{\text{SS}})$ be a (d, t) -LECSS with $E_{\text{SS}} : \{0, 1\}^{k'} \rightarrow \{0, 1\}^n$ and distance $d > n/4$. Let $(E_{\text{AMD}}, D_{\text{AMD}})$ be a ρ -secure AMD code where $E_{\text{AMD}} : \{0, 1\}^k \rightarrow \{0, 1\}^{n'}$ and $k' = n'$. Define the composed code (Enc, Dec) by*

$$\text{Enc}(s) \stackrel{\text{def}}{=} E_{\text{SS}}(E_{\text{AMD}}(s)), \quad \text{Dec}(c) \stackrel{\text{def}}{=} \begin{cases} \perp & \text{if } D_{\text{SS}}(c) = \perp \\ D_{\text{AMD}}(D_{\text{SS}}(c)) & \text{otherwise.} \end{cases}$$

Then (Enc, Dec) is non-malleable with respect to the family \mathcal{F}_{BIT} with error $\epsilon \leq \max(\rho, 2^{-\Omega(t)})$.

PROOF. We first sketch the high-level idea behind the proof. We must show that for each $f \in \mathcal{F}_{\text{BIT}}$, there is a distribution D_f which satisfies Definition 3.2. To do so, we look at how many of the component functions f_i that make up f are of the form “set to 0” or “set to 1”. If this number is too small (less than t), then tampering via f is really equivalent to adding some offset Δ to $E_{\text{AMD}}(s)$, from a distribution which is independent of $E_{\text{AMD}}(s)$. Therefore, there is some universal probability with which $\Delta = 0$ (in which case D_f should output same) and $\Delta \neq 0$ (in which case D_f should output \perp since the security of the AMD codes ensure that such modifications are likely to be detected). On the other hand, if the number of f_i that are “set to 0”/“set to 1” is too high (more than $n - t$), then the values in the positions of $\tilde{c} = f(\text{Enc}(s))$ that are not “set” are uniformly random. Therefore, D_f can sample the value of \tilde{c} independently of s . Lastly, we show that when the number of f_i that are “set to 0”/“set to 1” is in between t and $n - t$, then the erroneous codeword decodes to \perp with overwhelming probability. This is where we need the distance of the LECSS to be $d > n/4$.⁸ Essentially, we show that in this case there exists some codeword c^* such that the tampered value \tilde{c} is neither likely to match c^* exactly nor is it likely to be far away enough from it to be a valid codeword under the LECSS scheme.

Fix any tampering function $f = (f_1, \dots, f_n) \in \mathcal{F}$. We show that there exists a distribution D_f which satisfies the definition of non-malleability. To help with notation, for every $s \in \{0, 1\}^k$, we

⁷The construction of [28] defines an AMD code $(E_{\text{AMD}}, D_{\text{AMD}})$ which takes a message $s = (s_1, \dots, s_d) \in \mathbb{F}^d$ and outputs $E_{\text{AMD}}(s) = (s, x, f(x, s)) \in \mathbb{F}^{d+2}$ where $x \in \mathbb{F}$ is uniformly random and $f(x, s) = x^{d+2} + \sum_{i=1}^d s_i x^i$. The decoding $D_{\text{AMD}}(c)$ checks that a codeword c is of the form $c = (s', x', f(x', s'))$ for some s', x' and if so outputs s' . For any $\Delta = (\Delta_1, \Delta_2, \Delta_3)$ with $\Delta \neq 0$ the only way that $D_{\text{AMD}}(E_{\text{AMD}}(s) + \Delta) \neq \perp$ is if $f(x + \Delta_2, s + \Delta_1) = f(x, s) + \Delta_3$. By solving for x , it is easy to see that the above occurs only if x is a root of a non-zero polynomial of degree at most $d + 1$ which happens with probability at most $(d + 1)/|\mathbb{F}|$.

⁸Interestingly, the requirement that $d > n/4$ is not an artifact of the proof-technique but an optimal bound at this level of generality. In particular, there are examples of LECSS constructions with $d < n/4$ for which our construction would not be secure.

define the random variables:

$$C^s \stackrel{\text{def}}{=} \text{Enc}(s), \tilde{C}^s \stackrel{\text{def}}{=} f(C^s), \Delta^s \stackrel{\text{def}}{=} \tilde{C}^s - C^s, \tilde{S}^s \stackrel{\text{def}}{=} \text{Dec}(\tilde{C}^s).$$

We use the notation C_i^s ($\tilde{C}_i^s, \Delta_i^s$, respectively) to denote the bit of C^s (\tilde{C}^s, Δ^s , respectively) in position $i \in \{1, \dots, n\}$. We also define the distribution $\text{Patch}(D_f, s)$ which samples $\tilde{s} \leftarrow D_f$ and outputs s if $\tilde{s} = \text{same}$ and \tilde{s} otherwise. To prove the theorem, we must show that there exists a distribution D_f such that $\text{SD}(\tilde{S}^s, \text{Patch}(D_f, s)) \leq \varepsilon$ for all $s \in \{0, 1\}^k$. Let q be the number of positions $i \in \{1, \dots, n\}$ for which f_i is one of the functions “set to 0”/“set to 1.” We split the rest of the proof into four cases based on the value of q .

Case 1, $q \leq t$. In this case, the distribution of Δ^s is the same for all s . In particular:

- For i such that $f_i = \text{“keep”}$, $\Delta_i^s = 0$. For i such that $f_i = \text{“flip”}$, $\Delta_i^s = 1$.
- Let A be the set of i for which $f_i = \text{“set to 0”}$ or $f_i = \text{“set to 1”}$. For $i \in A$, Δ_i^s are uniformly random over $\{0, 1\}$ and independent.

Since $|A| = q \leq t$, we rely on the t -secrecy of the LECSS, which tells us that the random variables $\{C_i^s\}_{i \in A}$ are independent and uniform over $\{0, 1\}$. Since $\{f_i(C_i^s)\}_{i \in A}$ are constants, we see that $\{\Delta_i^s = C_i^s - f_i(C_i^s)\}_{i \in A}$ are uniformly random and independent variables for all s .

Therefore, there is a universal distribution Δ such that $\Delta \equiv \Delta^s$ for all s , and so

$$\begin{aligned} \tilde{S}^s &\equiv \text{Dec}(\tilde{C}^s) \equiv \text{D}_{\text{AMD}}(\text{D}_{\text{SS}}(C^s + \Delta^s)) \\ &\equiv \text{D}_{\text{AMD}}(\text{D}_{\text{SS}}(C^s) + \text{D}_{\text{SS}}(\Delta^s)) \equiv \text{D}_{\text{AMD}}(\text{E}_{\text{AMD}}(s) + \text{D}_{\text{SS}}(\Delta^s)) \\ &\equiv \text{D}_{\text{AMD}}(\text{E}_{\text{AMD}}(s) + \text{D}_{\text{SS}}(\Delta)), \end{aligned} \tag{3}$$

where Equation (3) follows by the linearity property of the LECSS. Therefore:

- (1) Conditioned on $\text{D}_{\text{SS}}(\Delta) \neq 0$, $\Pr[\tilde{S}^s = \perp] \geq (1 - \rho)$.
This follows by the security of the AMD code, and the fact that $\text{D}_{\text{SS}}(\Delta)$ is independent of the randomness of the AMD encoding.
- (2) Conditioned on $\text{D}_{\text{SS}}(\Delta) = 0$, $\Pr[\tilde{S}^s = s] = 1$.

We define the distribution D_f which samples $\delta \leftarrow \Delta$ and outputs same if $\text{D}_{\text{SS}}(\delta) = 0$ and \perp otherwise. By conditions (1) and (2) it is easy to see that, for all $s \in \{0, 1\}^k$,

$$\text{SD}(\tilde{S}^s, \text{Patch}(D_f, s)) \leq \rho,$$

which concludes the proof of Case 1.

Case 2, $q \geq n - t$. In this case, the distribution of \tilde{C}^s is the same for all s . In particular:

- For $f_i = \text{“set to 0”}$, $\tilde{C}_i^s = 0$. For $f_i = \text{“set to 1”}$, $\tilde{C}_i^s = 1$.
- Let B be the set of i for which $f_i = \text{“keep”}$ or $f_i = \text{“flip”}$. For $i \in B$, \tilde{C}_i^s are uniformly random over $\{0, 1\}$ and independent.

This is because $|B| = n - q \leq t$. By the t -secrecy of the LECSS, the random variables $\{C_i^s\}_{i \in B}$ are uniformly random and independent and hence so are the variables $\{\tilde{C}_i^s = f_i(C_i^s)\}_{i \in B}$.

Therefore, there is a universal distribution \tilde{C} such that $\tilde{C} \equiv \tilde{C}^s$ for all s and so $\tilde{S}^s \equiv \text{Dec}(\tilde{C}^s) \equiv \text{Dec}(\tilde{C})$. We define the distribution D_f which samples \tilde{C} as above, and computes $\text{Dec}(\tilde{C})$. It is clear that, for all $s \in \{0, 1\}^k$,

$$\text{SD}(\tilde{S}^s, \text{Patch}(D_f, s)) \leq \text{SD}(\tilde{S}^s, D_f) = 0,$$

which concludes the proof of Case 2.

Case 3, $t < q \leq n/2$. In this case, the distribution D_f just outputs \perp . Fix any $s \in \{0, 1\}^k$. It is easy to see that

$$\Pr[\tilde{S}^s \neq \perp] \leq \Pr[\text{Dec}(\tilde{C}^s) \neq \perp] \leq \Pr[D_{SS}(\Delta^s) \neq \perp].$$

We are now only left to show that $\Pr[D_{SS}(\Delta^s) \neq \perp]$ is small. Define A as the set of indices i for which f_i is of the form “set to 0”/“set to 1,” so that $|A| = q$. Note that $\{\Delta_i^s\}_{i \notin A}$ are completely fixed by f to some constants (Δ_i^s is fixed to 0 if f_i is “keep” and fixed to 1 if f_i is “flip”). Let $\delta^* \in \{0, 1\}^n$ be any value which is consistent with the fixed bits of Δ so that $\{\Delta_i^s = \delta_i^*\}_{i \notin A}$, and for which $D_{SS}(\delta^*) \neq \perp$. If no such value exists, then we are done since $D_{SS}(\Delta^s) = \perp$ with probability 1, so let us assume that some such value exists.

It is easy to see that $\Pr[\Delta^s = \delta^*] \leq \frac{1}{2^t}$ since the random variables $\{\Delta_i^s\}_{i \in A}$ are uniform over $\{0, 1\}$ and t -wise independent. So Δ^s is unlikely to match δ^* exactly. On the other hand, we show that the Hamming distance $d_H(\Delta^s, \delta^*)$ is unlikely to be too large. In particular, the expected value of this distance is

$$\mathbb{E}[d_H(\Delta^s, \delta^*)] = \mathbb{E}\left[\sum_{i=1}^n d_H(\Delta_i^s, \delta_i^*)\right] = \mathbb{E}\left[\sum_{i \in A} d_H(\Delta_i^s, \delta_i^*)\right] = \sum_{i \in A} \mathbb{E}[d_H(\Delta_i^s, \delta_i^*)] = \frac{q}{2},$$

where the second equality follows by the fact that $\Delta_i^s = \delta_i^*$ for $i \notin A$, the third equality follows by the linearity of expectation, and the fourth equality follows since each Δ_i^s for $i \in A$ is individually uniform. We now wish to bound the probability that $d_H(\Delta^s, \delta^*) = \sum_{i \in A} d_H(\Delta_i^s, \delta_i^*)$ is greater than the distance d of the LECSS scheme. Since $\{d_H(\Delta_i^s, \delta_i^*)\}_{i \in A}$ are t -wise independent random variables, we can use (Chernoff-Hoeffding type) tail inequalities for bounded independence (see [11, 65]) to bound the above probability. In particular,

$$\begin{aligned} \Pr[d_H(\Delta^s, \delta^*) \geq d] &\leq \Pr[|d_H(\Delta^s, \delta^*) - q/2| \geq d - q/2] \\ &\leq \left(\frac{nt}{(d - q/2)^2}\right)^{t/2} \leq \left(\frac{nt}{(d - n/4)^2}\right)^{t/2} \\ &\leq \left(\frac{t}{n(d/n - 1/4)^2}\right)^{t/2}, \end{aligned} \tag{4}$$

where Equation (4) follows by the tail inequality of Bellare and Rompel [11], Lemma 2.2. Finally, we see that

$$\Pr[D(\Delta^s) \neq \perp] \leq \Pr[\Delta^s = \delta^* \vee d_H(\Delta^s, \delta^*) \geq d] \leq \frac{1}{2^t} + \left(\frac{t}{n(d/n - 1/4)^2}\right)^{t/2},$$

which concludes the analysis of Case 3.

Case 4, $n/2 < q < n - t$. In this case, the distribution D_f just outputs \perp . Fix any $s \in \{0, 1\}^k$. It is easy to see that $\Pr[\tilde{S}^s \neq \perp] \leq \Pr[D_{SS}(\tilde{C}^s) \neq \perp]$ and so we are left to show that this last probability is small. To show this, we essentially re-use the analysis of Case 3 and note the symmetry between the distribution of \tilde{C}^s in this case and Δ^s in Case 3. In particular, define B as the set of indices i for which f_i = “keep”/“flip,” so $|B| = n - q$ and hence $t \leq |B| \leq n/2$. Note that $\{\tilde{C}_i^s\}_{i \notin B}$ are completely fixed by f . Let $\tilde{c}^* \in \{0, 1\}^n$ be any value which is consistent with the fixed portion of \tilde{C}^s so that $\{\tilde{C}_i^s = \tilde{c}_i^*\}_{i \notin B}$. Again, if no such value exists, then we are done. Otherwise, we re-use the exact same analysis as in Case 3 to argue that

$$\Pr[D_{SS}(\tilde{C}^s) \neq \perp] \leq \Pr[\tilde{C}^s = \tilde{c}^* \vee d_H(\tilde{C}^s, \tilde{c}^*) \geq d] \leq \frac{1}{2^t} + \left(\frac{t}{n(d/n - 1/4)^2}\right)^{t/2},$$

which concludes the analysis of Case 4. Since these are the only possible cases for q , this concludes the proof of the theorem. \square

We note that the idea of composing AMD codes together with linear secret-sharing schemes also appeared explicitly in [28] (and implicitly in [14, 60, 61]) in the context of “robust secret sharing,” but the application in this work and the main ideas behind our proof are very different.

4.2 Instantiating the Construction

To actually instantiate a construction of the non-malleable code defined above, we need constructions of AMD codes and of LECSS schemes. For the former, very efficient constructions of AMD codes were given by [28] where, to achieve security $\rho \leq 2^{-\lambda}$, the encoding has an additive overhead of roughly 2λ bits. On the other hand, LECSS schemes have not been explicitly defined or constructed in literature. We note that there is a clear connection between LECSS schemes, linear ramp secret-sharing schemes, and error-correcting codes. Ramp secret-sharing schemes have mostly been considered for larger alphabet sizes (e.g., Shamir Secret Sharing [66]) and not over bits; an exception is the work of Pueyo et al. [64]. However, no known constructions of ramp secret-sharing schemes over bits achieve distance $d > n/4$ over bits, as required for our application.

We now show how to instantiate LECSS schemes with large d, t efficiently, based on random linear error-correcting codes (recall that our construction does not require efficient error decoding). The main ideas for this construction are based on the work of Chen et al. [21], which shows how to relate the secrecy t of a LECSS to the dual distance of an underlying error-correcting code. In addition, it shows that a random linear error-correcting code will (with overwhelming probability) simultaneously achieve a large distance d and have a sub-code with large dual distance yielding a large t . Since our terminology and exact construction differ slightly, we state and prove the following lemma, which is essentially equivalent to the result of [21].

LEMMA 4.2 (BASED ON [21]). *Let G be a $(k + \ell) \times n$ generator matrix over some field \mathbb{F} , generating a code $C \subseteq \mathbb{F}^n$ with distance d . Let \hat{G} be the sub-matrix consisting of the last ℓ rows of G , let $\hat{C} \subseteq C$ be the sub-code generated by \hat{G} , and let \hat{C}^\perp be the dual code to \hat{C} . Let \hat{d}^\perp be the distance of \hat{C}^\perp .*

We define the coding scheme (E_{SS}, D_{SS}) , with source messages $s \in \mathbb{F}^k$. The encoding algorithm computes $E_{SS}(s) \stackrel{\text{def}}{=} (s, r)G$ for a uniformly random $r \in \mathbb{F}^\ell$. The decoding algorithm $\text{Dec}(c)$ checks if c is a codeword in C : if so, it decodes c to $(s', r') \in \mathbb{F}^k \times \mathbb{F}^\ell$ and outputs s' ; otherwise, it outputs \perp . Then (E_{SS}, D_{SS}) is a (d, t) -LECSS scheme where $t = \hat{d}^\perp - 1$.

PROOF. The linearity and distance d of the LECSS follows directly from the linearity and distance of C . To analyze privacy, note that for any $s \in \mathbb{F}^k$, we have $E_{SS}(s) = (s, r)G = (s, 0)G + r\hat{G}$ and so we just need to show that the components of $\hat{c} = r\hat{G}$ are individually uniform and t -wise independent for a random r . To show this, we use a well-known fact from coding theory that, if the distance of \hat{C}^\perp is \hat{d}^\perp , then any $t = \hat{d}^\perp - 1$ columns of \hat{G} must be linearly independent. Therefore, for any fixed set of indices $A \subseteq \{1, \dots, n\}$ of size $|A| = t$, multiplication by \hat{G} is a surjective linear function when restricted to the columns $i \in A$ and so the components of $\hat{c} = r\hat{G}$ in positions $i \in A$ are uniformly random and independent. \square

Unfortunately, it is unclear how to construct an error-correcting code C which has a good distance $d > n/4$ and having a sub-code \hat{C} with large dual distance \hat{d}^\perp over a binary alphabet (even if we are not concerned about simultaneously achieving a good rate). We do not know of any explicit constructions of such codes, and it remains as an interesting open problem to find one. However, it is relatively easy to show that a *random linear error-correcting code* is likely (with overwhelming probability) to simultaneously achieve very high distance d for C and high dual distance \hat{d}^\perp for \hat{C}^\perp .

as defined in Lemma 4.2. We note that, in our context, using random linear error-correcting codes is *very efficient*, since the representation size of such code, the encoding procedure, checking if a value c is a codeword, and decoding a codeword *without errors* are all very efficient. In particular, the construction in Lemma 4.2 does *not* require the ability to efficiently *correct* errors, which is believed to be hard for random linear codes. The following result, explicitly shown by Chen et al. [21] (see Theorem 4 and Corollary 1), follows from a Gilbert-Varshamov type of argument showing that random linear error-correcting codes achieve good distance and also have sub-codes with good dual distance.

LEMMA 4.3 ([21]). *Let G be a uniformly random generator matrix from the space of $(k + \ell) \times n$ matrices over \mathbb{F}_2 with linearly independent rows. Let $C, \hat{C}, \hat{C}^\perp, d$, and \hat{d}^\perp be defined as in Lemma 4.2. Then for any $0 < \alpha, \hat{\alpha} < \frac{1}{2}$,*

$$\Pr[(d > \alpha n) \text{ and } (\hat{d}^\perp > \hat{\alpha} n)] > 1 - 2^{k+\ell-n(1-H(\alpha))} - 2^{\ell-nH(\hat{\alpha})},$$

where the probability is taken (only) over the choice of G , and $H(x) \stackrel{\text{def}}{=} -x \log(x) - (1-x) \log(1-x)$.

Using the above two lemmas, we are now ready to prove the following.

THEOREM 4.4. *For any constant $\delta > 0$, any $n \in \mathbb{N}$, there exist efficient non-malleable codes with respect to the family \mathcal{F}_{BIT} , with codeword size n , message size $k \geq (1 - H(1/4) - \delta)n$, and security $\varepsilon = 2^{-\Omega(n)}$, where $H(x) \stackrel{\text{def}}{=} -x \log(x) - (1-x) \log(1-x)$ is the Shannon entropy function. In particular, the rate of the code can be made to approach $k/n \approx (1 - H(1/4)) \approx .1887$. Moreover, there is an efficient procedure which, given k and n , outputs a description of such a code with probability $1 - 2^{-\Omega(n)}$.*

PROOF. First, given $\delta > 0$, there are constants $\delta_0 > 0, \delta_1 > 0$ such that

$$(1 - H(1/4) - \delta) = (1 - H(1/4 + \delta_0) - \delta_1).$$

Let $k' = (1 - H(1/4 + \delta_0) - \delta_1/4)n$ and $\ell = (\delta_1/4)n$. Let $\hat{\alpha}$ be a constant such that $H(\hat{\alpha}) \leq \delta_1/8$. If we choose a random $(k' + \ell) \times n$ matrix G over \mathbb{F}_2 , and define $C, \hat{C}, \hat{C}^\perp, d, \hat{d}^\perp$ as in Lemma 4.2, then using the bounds of Lemma 4.3, we will get $d \geq (1/4 + \delta_0)n$ and $\hat{d}^\perp \geq \hat{\alpha}n$ with probability at least

$$1 - 2^{k'+\ell-n(1-H(1/4+\delta_0))} - 2^{\ell-nH(\hat{\alpha})} = 1 - 2^{-(\delta_1/2)n} - 2^{-(\delta_1/8)n} = 1 - 2^{-\Omega(n)}.$$

Assuming this is the case, we now show how to instantiate our non-malleable code using G , so that we achieve the claimed parameters.

- (1) Using Lemma 4.2, we see that G gives rise to an efficient (d, t) -LECSS scheme with message size k' and codeword size n , and with $d \geq (1/4 + \delta_0)n$ and $t = \hat{d}^\perp - 1 \geq \delta_2 n$ for some constant $\delta_2 < \delta_0^2/2$.
- (2) Using the result of [28], there exist AMD codes with message size k , codeword size k' , and security

$$\rho \leq \frac{k}{2^{(k'-k)/2}} \leq \frac{n}{2^{(\delta_1/4)n}} = 2^{-\Omega(n)}.$$

Now we use Theorem 4.1 to combine (1) and (2) above, yielding a non-malleable code with message size k , codeword size n , and security

$$\varepsilon \leq \max \left(\rho, \frac{1}{2^{\delta_2 n}} + \left(\frac{\delta_2 n}{n(\delta_0)^2} \right)^{(\delta_2 n)/2} \right) \leq 2^{-\Omega(n)}. \quad \square$$

4.3 Improving Efficiency

One issue with our construction is that the rate k/n of the scheme is less than $1/5$ and therefore relatively poor. We now address this issue by showing that *any* non-malleable code with respect to \mathcal{F}_{BIT} , with any *constant* rate k/n , can be converted into *computationally secure* non-malleable code with respect to \mathcal{F}_{BIT} , where the new rate k'/n' asymptotically approaches 1. The main idea is to combine an encoding scheme (Enc, Dec) from the previous sections with a *symmetric-key authenticated encryption scheme* (AuthEncrypt, AuthDecrypt) in the following way: to encode a message s we take a random key key and store $(\text{Enc}(\text{key}), \text{AuthEncrypt}_{\text{key}}(s))$. The decoding of (c_0, c_1) is straightforward: first compute $\text{key}' = \text{Dec}(c_0)$, and then output $\text{AuthDecrypt}_{\text{key}'}(c_1)$.

We construct our symmetric-key authenticated-encryption scheme using an encryption scheme (Encrypt, Decrypt) with key size k_0 , and a message authentication code (MAC) (Tag, Verify) with key size k_1 . The encryption scheme (Encrypt, Decrypt) has to always satisfy the following: $\text{Decrypt}(\text{key}_0, \text{Encrypt}(\text{key}_0, m)) = m$. Security of (Encrypt, Decrypt) is defined as follows: for every two messages m and m' of equal length, and a random key_0 , we have that $\text{Encrypt}(\text{key}_0, m)$ and $\text{Encrypt}(\text{key}_0, m')$ are computationally indistinguishable.

The authentication scheme (Tag, Verify) has always to satisfy: $\text{Verify}(\text{key}_1, m, \text{Tag}(\text{key}_1, m)) = \text{yes}$. Security of (Tag, Verify) is defined as follows: for every message m and for a random key_1 no poly-time adversary \mathcal{A} , after seeing $(m, \text{Tag}(\text{key}_1, m))$, can output (t, m') such that $\text{Verify}(\text{key}_1, m', t) = \text{yes}$ and $m \neq m'$, except with a negligible probability.

An authenticated encryption scheme that we use has a key (k_0, k_1) . On a message m it outputs $\text{Tag}(k_1, \text{Encrypt}(k_0, m))$ (cf. [9] for a general discussion on how to combine authentication with encryption).

The following theorem shows that, assuming the existence of one-way functions, there are efficient *computational* non-malleable codes with respect to the family \mathcal{F}_{BIT} , where the rate k/n asymptotically approaches 1.

THEOREM 4.5. *Assume that (Enc, Dec) is a non-malleable code with respect to the family \mathcal{F} of bitwise independent tampering functions with message-size k and codeword-size n . Let (Encrypt, Decrypt) and (Tag, Verify) be as above. Define*

$$\text{Enc}'(s) \stackrel{\text{def}}{=} \left\{ \begin{array}{l} (\text{key}_0, \text{key}_1) \leftarrow U_{k_0+k_1}, x \leftarrow \text{Encrypt}(\text{key}_0, s), t \leftarrow \text{Tag}(\text{key}_1, x), \\ c \leftarrow \text{Enc}(\text{key}_0 || \text{key}_1) \\ \text{Output: } c = (c, x, t) \end{array} \right\},$$

$$\text{Dec}'(c, x, t) \stackrel{\text{def}}{=} \left\{ \begin{array}{l} \tilde{\text{key}}_0 || \tilde{\text{key}}_1 = \text{Dec}(c), \\ \text{If } \tilde{\text{key}}_0 || \tilde{\text{key}}_1 = \perp \text{ or } \text{Verify}(\text{key}_1, x, t) \neq \text{yes output } \perp, \\ \text{else output } \text{Decrypt}(\text{key}_0, x) \end{array} \right\}.$$

Then $(\text{Enc}', \text{Dec}')$ is a computational non-malleable code with respect to \mathcal{F} .

PROOF. Intuitively, the reason why the theorem holds is that the adversary has a choice to either (1) make the key encoding $(\text{key}_0, \text{key}_1)$ decode to \perp —in this case the whole encoding decodes to \perp not matter what the encoded message was; or (2) not change the encoded key $(\text{key}_0, \text{key}_1)$ —in this case the adversary cannot substitute the encode message s with any other s' without breaking the MAC; moreover, the security of the encryption scheme guarantees the probability that the decoding yields \perp has to be (almost) the same for every message s ; finally, he can (3) change the encoded key to some random key $(\tilde{\text{key}}_0, \tilde{\text{key}}_1)$ (that has a distribution that is independent on $(\text{key}_0, \text{key}_1)$)—in this case it follows from the security of the encryption scheme that the decoded message cannot depend on the encoded message s , as otherwise one could construct a distinguisher that breaks the encryption scheme by sampling $(\tilde{\text{key}}_0, \tilde{\text{key}}_1)$ himself.

Let f' be a function that is tampering the output of $(\text{Enc}', \text{Dec}')$. We will construct a distribution $D'_{f'}$ such that for every s (1) is satisfied. Since the output of Enc' has a form (c, x, t) , and since f' only modifies the individual bits, we can think of f' as of three functions f'_0 , f'_1 , and f'_2 , where $f'(c, x, t) = (f'_0(c), f'_1(x), f'_2(t))$. The distribution $D'_{f'}$ is constructed as follows. First, we take an arbitrary message s and a random pair $(\text{key}_0, \text{key}_1)$. Then, we let $x^s \leftarrow \text{Encrypt}(\text{key}_0, s)$ and $t^s \leftarrow \text{Tag}(\text{key}_1, x)$. Form the assumption that (Enc, Dec) is a non-malleable code there exists a distribution $D_{f'_0}$ such that (1) holds. We sample \tilde{c} from this distribution. If $\tilde{c} = \text{same}$, we let $(\tilde{\text{key}}_0, \tilde{\text{key}}_1) := (\text{key}_0, \text{key}_1)$, and otherwise we let $(\tilde{\text{key}}_0, \tilde{\text{key}}_1) := \tilde{c}$. We then output

$$\text{out}^s := \begin{cases} \perp & \text{if } \text{Verify}(\tilde{\text{key}}_1, f'_1(x^s), f'_2(t^s)) \neq \text{yes or } \tilde{c} = \perp, \\ \text{same} & \text{if } f'_1(x^s) = x^s \text{ and } \text{Verify}(\text{key}_1, f'_1(x^s), f'_2(t^s)) = \text{yes} \\ & \text{and } \tilde{c} = \text{same}, \\ \text{Dec}(\tilde{\text{key}}_0, f'_1(x^s)) & \text{otherwise.} \end{cases}$$

To prove that this construction is correct, we need to show that for every s^0 and s^1 the distributions out^{s^0} and out^{s^1} are computationally indistinguishable. After showing this we will be done, since this will mean that the result of our simulation (with an arbitrary value s) is computationally indistinguishable from a result of the simulation with any other s , and out^s with same substituted with s is simply distributed identically to the output of the experiment of the left-hand side of Equation (1). Obviously, the distribution of \tilde{c} does not depend on s . Therefore, we can consider separately the following cases.

– $\tilde{c} = \perp$: In this case, obviously out^s is always equal to \perp , no matter what s is.

– $\tilde{c} = \text{same}$: In this case, first observe that

$$\left| \Pr [\text{out}^{s^0} = \perp] - \Pr [\text{out}^{s^1} = \perp] \right| \text{ is negligible.}$$

This is because otherwise one could construct an adversary that breaks the security of the encryption scheme $(\text{Encrypt}, \text{Decrypt})$ by distinguishing if x is an encryption of s^0 and s^1 in the following way: generate a random key_1 and output 1 if and only if $\text{Verify}(\text{key}_1, f'_1(x), f'_2(\text{Tag}(\text{key}_1, x)))$.

On the other hand, if $\text{out}^s \neq \perp$ and $f'_1(x^s) \neq x^s$, then we can break the security of the MAC in the following way: after seeing a tag t^s on x^s an adversary can output a message $f'_1(x^s)$ with a tag $f'_1(x^s)$.

– $\tilde{c} = (\tilde{\text{key}}_0, \tilde{\text{key}}_1)$: If in this case out^{s^0} can be distinguished from out^{s^1} by some poly-time machine \mathcal{A} , then one can easily distinguish if x is an encryption of s^0 and s^1 in the following way.

- (1) if $\text{Verify}(\tilde{\text{key}}_1, f'_1(x), f'_2(\text{Tag}(\tilde{\text{key}}_1, x))) = \perp$ then output \perp ,
- (2) otherwise output $\mathcal{A}(\text{Dec}(\tilde{\text{key}}_0, f'_1(x)))$.

This is because for both $i \in \{0, 1\}$ we have $\text{out}^{s^i} = \text{Dec}(\tilde{\text{key}}_0, f'_1(x^{s^i}))$. □

5 A GENERAL RESULT FOR TAMPERING-FUNCTION FAMILIES OF BOUNDED SIZE

5.1 A Probabilistic Method Approach

We now show that, for *any* “small enough” function family \mathcal{F} , there *exists* a (possibly inefficient) coding scheme which is non-malleable with respect to \mathcal{F} . Moreover, we show that for a fixed “small enough” function family \mathcal{F} , a *random* coding scheme is likely to be non-malleable with respect to \mathcal{F} with overwhelming probability. Unfortunately, random coding schemes cannot be efficiently represented, nor is the encoding/decoding function likely to be efficient. Therefore, this

result should merely be thought of as showing “possibility” and providing a target that we should then strive to match constructively. Moreover, this result also highlights the difference between “error correction/detection” and “non-malleability” since a result of this form could not be true for the former notions.

Let us now clarify what we mean by “small enough.” The family \mathcal{F}_{all} of all functions $f : \{0, 1\}^n \rightarrow \{0, 1\}^n$ has size $\log(\log(|\mathcal{F}_{\text{all}}|)) = n + \log(n)$. Clearly, no code is non-malleable with respect to the family \mathcal{F}_{all} , since this family includes, for example, the function that decodes the codeword, flips the last bit of the source message, and re-encodes it. However, we show that for any function family \mathcal{F} of slightly smaller size $\log(\log(|\mathcal{F}|)) < n$, there exist codes that are non-malleable with respect to \mathcal{F} .

Lastly, let us clarify what we mean by “random” coding scheme. We can define a coding scheme completely by only specifying the decoding function $\text{Dec} : \{0, 1\}^n \rightarrow \{0, 1\}^k \cup \perp$, which then implicitly defines an encoding function Enc that maps a source message s uniformly at random into the set $\{c \in \{0, 1\}^n \mid \text{Dec}(c) = s\}$. For us, a random coding scheme is a scheme of this type, where the decoding function is chosen *uniformly* at random from all functions $\text{Dec} : \{0, 1\}^n \rightarrow \{0, 1\}^k$; that is, we specify the decoding function by taking each value $c \in \{0, 1\}^n$ and letting it decode to a uniformly random source message $s \in \{0, 1\}^k$. Note that there are *no* invalid values $c \in \{0, 1\}^n$ that decode to \perp .

THEOREM 5.1. *Let \mathcal{F} be any function family consisting of functions $f : \{0, 1\}^n \rightarrow \{0, 1\}^n$. Let $\varepsilon, \rho > 0$ be arbitrary values and $k, n > 0$ be integers such that*

$$\log(\log(|\mathcal{F}|)) > 3k + \log(k) - n + 2 \log(1/\varepsilon) + \log(\log(1/\rho)) + 9. \quad (5)$$

Then there exists a strongly non-malleable code with respect to \mathcal{F} , with k -bit source messages and n -bit codewords, and exact security ε . Moreover, a uniformly random decoding function $\text{Dec} : \{0, 1\}^n \rightarrow \{0, 1\}^k$ gives rise to such a code with probability $\geq 1 - \rho$.

PROOF. Let $K = 2^k, N = 2^n$. It will be useful for us to think of a function $f : \{0, 1\}^n \rightarrow \{0, 1\}^n$ as a directed graph $G_f = (V, E_f)$ with vertex set $V = \{0, 1\}^n$ and edges $E_f = \{(v, \tilde{v}) \mid v \in V, f(v) = \tilde{v}\}$. Notice that the out-degree is 1 and that there may be some self-loops in G . Notice that a randomly chosen decoding function, essentially labels each vertex $v \in V$ with a uniformly random value $s \in \{0, 1\}^k$. For any $s, \tilde{s} \in \{0, 1\}^k$, we define the following random variables (over the choice of the decoding function Dec):

- Let $V(s)$ be the set of vertices $v \in V = \{0, 1\}^n$ such that $\text{Dec}(v) = s$.
- Let $E_f(s \rightarrow \tilde{s})$ be the set of edges (v, \tilde{v}) where $v \neq \tilde{v}$ and $\text{Dec}(v) = s, \text{Dec}(\tilde{v}) = \tilde{s}$.
- Let $E_f(s \rightarrow \text{same})$ be the set of self-loop edges (v, v) for which $\text{Dec}(v) = s$.
- Let $E_f(* \rightarrow \tilde{s}) \stackrel{\text{def}}{=} \bigcup_{s \in \{0, 1\}^k} E_f(s \rightarrow \tilde{s})$ (we will also use this notation for $\tilde{s} = \text{same}$).
- For $A \subseteq V \cup \{\text{same}\}$, define $E_f(s \rightarrow A) \stackrel{\text{def}}{=} \bigcup_{\tilde{s} \in A} E_f(s \rightarrow \tilde{s})$, and $E_f(* \rightarrow A) \stackrel{\text{def}}{=} \bigcup_{\tilde{s} \in A} E_f(* \rightarrow \tilde{s})$.

We use these variables to define a density function $D_f : \{0, 1\}^k \cup \{\text{same}\} \rightarrow [0, 1]$ by $D_f(s) = |E_f(* \rightarrow s)|/N$. Then, for any function f and any $s \in \{0, 1\}^k$, we have

$$\text{SD}(\text{StrongNM}_s^f, D_f) = \max_{A \subseteq \{0, 1\}^k \cup \{\text{same}\}} \frac{|E_f[s \rightarrow A]|}{|V(s)|} - \frac{|E_f(* \rightarrow A)|}{N}. \quad (6)$$

Therefore, taking probabilities (only) over the choice of the decoding function Dec , let \mathcal{E} be the event that (Enc, Dec) is not an ε -secure non-malleable code with respect to \mathcal{F} . Then,

$$\Pr[\mathcal{E}] \tag{7}$$

$$\begin{aligned} &\leq \Pr[\exists f \in \mathcal{F}, s \in \{0, 1\}^k \text{ s.t. } \text{SD}(\text{StrongNM}_s^f, D_f) > \varepsilon] \\ &= \Pr\left[\exists f \in \mathcal{F}, s \in \{0, 1\}^k, A \subseteq \{0, 1\}^k \cup \{\text{same}\} \text{ s.t. } \frac{|E_f[s \rightarrow A]|}{|V(s)|} - \frac{|E_f[* \rightarrow A]|}{N} > \varepsilon\right] \\ &\leq \sum_{f \in \mathcal{F}} \sum_{s \in \{0, 1\}^k} \sum_{A \subseteq \{0, 1\}^k \cup \{\text{same}\}} \Pr\left[\frac{|E_f[s \rightarrow A]|}{|V(s)|} - \frac{|E_f[* \rightarrow A]|}{N} > \varepsilon\right]. \end{aligned} \tag{8}$$

We now fix some $f \in \mathcal{F}$, $s \in \{0, 1\}^k$, $A \subseteq \{0, 1\}^k \cup \{\text{same}\}$ and our goal is to upper bound:

$$\Pr\left[\frac{|E_f[s \rightarrow A]|}{|V(s)|} - \frac{|E_f[* \rightarrow A]|}{N} > \varepsilon\right]. \tag{9}$$

Let $\delta = \varepsilon/4$, $t = (\varepsilon/4)(N/K)$ and assume that

$$|V(s)| \geq (1 - \delta) \frac{N}{K}, |E_f(s \rightarrow A)| - \frac{|E_f[* \rightarrow A]|}{K} \leq t \tag{10}$$

then

$$\begin{aligned} \frac{|E_f[s \rightarrow A]|}{|V(s)|} - \frac{|E_f[* \rightarrow A]|}{N} &\leq \frac{|E_f[* \rightarrow A]|/K + t}{(1 - \delta)N/K} - \frac{|E_f[* \rightarrow A]|}{N} \\ &\leq \frac{\delta|E_f[* \rightarrow A]| + tK}{N(1 - \delta)} \\ &\leq 2\delta + 2tK/N \leq \varepsilon, \end{aligned}$$

where the inequality on the last line follows since $|E_f[* \rightarrow A]| \leq N$. Therefore, we show

$$(9) \leq \Pr[(10) \text{ does not hold}] \tag{11}$$

$$\begin{aligned} &\leq \Pr\left[|V(s)| < (1 - \delta) \frac{N}{K}\right] + \Pr\left[|E_f(s \rightarrow A)| - \frac{|E_f[* \rightarrow A]|}{K} > t\right] \\ &\leq e^{-\delta^2 \frac{N}{2K}} + 2e^{-t^2/8N} \\ &\leq e^{-\frac{\varepsilon^2 N}{32K}} + 2e^{-\frac{\varepsilon^2 N}{128K^2}} \leq 3e^{-\frac{\varepsilon^2 N}{128K^2}}, \end{aligned} \tag{12}$$

where we rely on the bounds from Lemma 5.2 and Lemma 5.3 (presented after this proof) for Equation (12). Now, continuing from Equation (8), we have

$$\begin{aligned} \Pr[\mathcal{E}] &\leq \sum_{f \in \mathcal{F}} \sum_{s \in \{0, 1\}^k} \sum_{A \subseteq \{0, 1\}^k \cup \{\text{same}\}} 3e^{-\frac{\varepsilon^2 N}{128K^2}} \\ &\leq 2^{\log(|\mathcal{F}|) + k + K + 3 - \frac{\varepsilon^2 N}{128K^2}}. \end{aligned} \tag{13}$$

We now bound Equation (13) by ρ by setting

$$\begin{aligned} \Pr[\mathcal{E}] \leq \rho &\quad \text{follows if} \quad \log(|\mathcal{F}|) + k + K + 3 + \log(1/\rho) < \varepsilon^2 N / 128K^2 \\ &\quad \text{follows if} \quad 128K^2 / \varepsilon^2 (\log(|\mathcal{F}|) + k + K + 3 + \log(1/\rho)) < N \\ &\quad \text{follows if} \quad n > \log(\log(|\mathcal{F}|)) + 3k + \log(k) + 2 \log\left(\frac{1}{\varepsilon}\right) + \log\left(\log\left(\frac{1}{\rho}\right)\right) + 9, \end{aligned}$$

where the last inequality is the assumption of our theorem. Therefore, with probability $1 - \rho$, \mathcal{E} does not occur and the random function Dec gives rise to a non-malleable code with respect to \mathcal{F} with security ε . In particular, some such codes with the above bound exist for any $\rho < 1$. This concludes the proof, up to the bounds from Lemma 5.2 and Lemma 5.3, which are shown next. \square

LEMMA 5.2. For any $\delta > 0$, $\Pr[|V(s)| < (1 - \delta)\frac{N}{K}] \leq e^{-\delta^2 \frac{N}{2K}}$.

PROOF. Follows directly by the (multiplicative) Chernoff bound, since each vertex $v \in V$ is labeled with a uniformly random value from $\{0, 1\}^k$. \square

LEMMA 5.3. Let $W = |E_f(s \rightarrow A)| - \frac{|E_f(* \rightarrow A)|}{K}$. Then, for any $t \geq 0$, $\Pr[W \geq t] \leq 2e^{-t^2/8N}$.

PROOF. For each edge $e = (v, \tilde{v})$, we define a random variable

$$C_e = \begin{cases} 0 & \text{if } e \notin E_f(* \rightarrow A) \\ 1 - 1/K & \text{if } e \in E_f(s \rightarrow A) \\ (-1/K) & \text{otherwise (if } e \in E_f(* \rightarrow A) \setminus E_f(s \rightarrow A)). \end{cases} \quad (14)$$

Then, $W = \sum_{e \in E} C_e$. We wish to prove that W is closely centered around 0. Unfortunately, the variables C_e are not independent. Our analysis is broken up into three claims. First we show that in any subset $U \subseteq E$ which does not contain *non-trivial cycles* (cycles with more than one edge), the sum of C_e over U forms a martingale. Second, we show that E can be partitioned into two such subsets $E = U_1 \cup U_2$. Lastly, we show that these two properties imply that W is closely centered at 0, using Azuma's inequality on martingales.

CLAIM 5.1. For any subset $U \subseteq E$ of edges which does not contain a non-trivial cycle (i.e., a cycle consisting of more than one edge), there is an ordering $e_1, \dots, e_{|U|}$ of the edges in U such that the random variables $W_i = \sum_{j=1}^i C_{e_j}$ form a Martingale.

PROOF. We can define a topological sort $v_1, \dots, v_{|V|}$ on the vertices V so that, for all edges $e = (v_i, v_j) \in U$, $i \geq j$. We can extend the ordering on vertices to one on edges by identifying each edge with its starting vertex (this is unique since the out-degree is 1). Think of the function Dec as being chosen by assigning the values $\text{Dec}(v_1), \text{Dec}(v_2), \dots, \text{Dec}(v_{|V|})$ in order, uniformly at random. Let $e_k = (v_i, v_j)$ be some edge in U . Then, the choice of $\text{Dec}(v_1), \dots, \text{Dec}(v_{i-1})$ completely determines the values $C_{e_1}, \dots, C_{e_{k-1}}$ and hence also the values of W_1, \dots, W_{k-1} . We consider two cases.

$v_i = v_j$: In this case, if $\text{same} \notin A$, then C_{e_k} is always 0. If $\text{same} \in A$, then e_k is always in $E_f(* \rightarrow A)$, and is also in $E_f(s \rightarrow A)$ iff $\text{Dec}(v_i)$ is labeled with s , which occurs with probability $1/K$. Therefore, $E[C_{e_k} | C_{e_1}, \dots, C_{e_{k-1}}] = 0$.

$v_i \neq v_j$: In this case, the value $\text{Dec}(v_j)$ is completely determined at this point since $i > j$. If $\text{Dec}(v_j) \notin A$, then C_{e_k} is always 0. Otherwise, it takes on the value $1 - 1/K$ with probability $1/K$ (the probability the $\text{Dec}(v_i) = s$) and $-1/K$ with probability $(1 - 1/K)$. Therefore, $E[C_{e_k} | C_{e_1}, \dots, C_{e_{k-1}}] = 0$.

So, in both cases $E[C_{e_k} | C_{e_1}, \dots, C_{e_{k-1}}] = 0$ and $E[W_k | W_1, \dots, W_{k-1}] = W_{k-1}$ as we wanted to show. \square

CLAIM 5.2. For any directed graph $G = (V, E)$ where each vertex has out-degree 1, there is a partition of E into two components U_1, U_2 such that neither component contains a non-trivial cycle.

PROOF. Since the out-degree is 1, each edge participates in at most one non-trivial cycle. Therefore, we can break up each non-trivial cycle separately, by placing half the edges into U_1 and the other half into U_2 . \square

(Continuing the proof of Lemma 5.3). Let U_1, U_2 be a partition of E as in Claim 5.2. Let $W_1 = \sum_{e \in U_1} C_e$, $W_2 = \sum_{e \in U_2} C_e$ so that $W = W_1 + W_2$. Then, by Claim 5.1 we have the martingales $W_1^{(1)}, W_1^{(2)}, \dots, W_1^{(|U_1|)} = W_1$ and $W_2^{(1)}, W_2^{(2)}, \dots, W_2^{(|U_2|)} = W_2$. So

$$\Pr[W \geq t] \leq \Pr[W_1 \geq t/2] + \Pr[W_2 \geq t/2] \leq e^{-t^2/8|U_1|} + e^{-t^2/8|U_2|} \leq 2e^{-t^2/8N},$$

where we use Azuma's inequality to bound the two probabilities (and note that $|U_1|, |U_2| \leq N$). \square

5.2 Constructions in the Random Oracle Model

It is not clear what the bound from Theorem 5.1 actually implies. For example, it does tell us that non-malleable codes exist with respect to all efficient functions, but this is misleading as we know that *efficient* non-malleable codes (and ultimately we are only interested in such) cannot be non-malleable with respect to this class.

Nonetheless, as we will explain below, the result by the probabilistic method does give us codes which are non-malleable with respect to very general classes of functions in the random oracle model. In particular, we can take a probabilistic method construction of non-malleable codes and convert it into a construction in the random oracle mode. There are two possible methods to do this.

By Theorem 5.1, a random decoding $\text{Dec} : \{0, 1\}^n \rightarrow \{0, 1\}^k$ will be non-malleable, thus we could instantiate Dec with a random oracle, but then it is not obvious how to implement the corresponding encoding Enc efficiently. We can fix this by relying on a random permutation instead of random oracle. Coron et al. [27] show how to construct a permutation \mathcal{P}^O (which can be queried from both sides) from a random oracle O so that \mathcal{P}^O is *indifferentiable* from a uniformly random permutation. Thus, instead of a random oracle, we can assume (cf. Proposition 2 in [55]) that we have oracle access to a uniformly random permutation \mathcal{P} and its inverse \mathcal{P}^{-1} over $\{0, 1\}^n$. We then define the encoding $\text{Enc}(x; r) \stackrel{\text{def}}{=} \mathcal{P}(x, r)$ for messages $x \in \{0, 1\}^k$ and randomness $r \in \{0, 1\}^{n-k}$. The decoding of $c \in \{0, 1\}^n$ are simply the first k bits of $\mathcal{P}^{-1}(c)$. The distribution on Dec is not exactly the distribution of a random decoding function as analyzed in Theorem 5.1, as now we are using a random permutation instead of a random function. However, since this does not change the probabilities significantly, it would be easy to modify the proof of Theorem 5.1 to make it work with a random permutation in place of a random function. In addition to needing a modified proof, this approach also requires us to bound the number of oracle queries made by the adversary in order to rely on the indistinguishability proof of [27].

Alternatively, we can avoid the above issues altogether by relying on follow-up works of [43, 52] which give probabilistic method constructions of non-malleable codes analogous to Theorem 5.1 (with even better parameters) but where both the encoding and decoding functions only make black-box use of a truly random function. Therefore, we can easily replace the random function by a random oracle in these constructions to get a construction of non-malleable codes in the random oracle model.

Thus, we can get a non-malleable code with respect to any \mathcal{F} as in Theorem 5.1 relative to a random oracle, and security holds even against an unbounded attacker that queries the oracle on all points. However, it is important to note that here the tampering functions $f \in \mathcal{F}$ cannot have access to the random oracle O .⁹ For this reason, this should *not* be considered a result in the random oracle model, where one usually assumes that the oracle is accessible by all parties.¹⁰ To get a result in the random oracle model, we must give all parties, *including the tampering functions*, access to the random oracle. Formally, codes that are strongly non-malleable in the random oracle model are defined as follows (the standard, *non-strongly* non-malleable codes, are defined analogously).

Definition 5.1 (Strong Non-Malleability in the Random Oracle Model). Let Enc and Dec be oracle machines, such that for every oracle O (with some fixed output length) the pair $(\text{Enc}^O, \text{Dec}^O)$ is a coding scheme. Let \mathcal{M} be a family of oracle machines (see Section 2) that take as inputs and

⁹Otherwise, for example, the class \mathcal{F} containing the single function f , where $f^O(c)$ decodes c and then re-encodes the message with the last bit flipped, is a counterexample.

¹⁰Security proofs in the random oracle model only exploit the fact that an exponential amount of uniform randomness is efficiently accessible by all parties, not that some entities (like, in our case, tampering functions) cannot access this randomness.

produce as outputs codewords of (Enc, Dec) . We say that (Enc, Dec) is *strongly non-malleable in the random oracle model with respect to tampering machines* \mathcal{M} if for any $s_0, s_1 \in \{0, 1\}^k$ and any $M \in \mathcal{M}$, we have

$$\text{StrongNM}_{s_0}^M \approx \text{StrongNM}_{s_1}^M,$$

$$\text{where } \text{StrongNM}_s^M \stackrel{\text{def}}{=} \left\{ \begin{array}{l} c \leftarrow \text{Enc}^O(s), \tilde{c} \leftarrow M^O(c), \tilde{s} = \text{Dec}^O(\tilde{c}) \\ \text{Output } \underline{\text{same}} \text{ if } \tilde{c} = c, \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\}$$

In the above experiment, O is an oracle containing a uniformly random function, and “ \approx ” can refer to statistical or computational indistinguishability. In the case of statistical indistinguishability, we say the scheme has exact security ε , if the statistical distance is at most ε .

We now define the following class of tampering machines.

Definition 5.2 (Closed Class of Tampering Machines). Consider a class \mathcal{M} of oracle machines. We say \mathcal{M} is *closed*, if for every oracle O and every $M \in \mathcal{M}$, there exists an $M' \in \mathcal{M}$, that never queries O , such that M' computes the function M^O .

Clearly, if \mathcal{M} is closed, then the subset $\hat{\mathcal{M}} \subset \mathcal{M}$ of machines that make no oracle queries constitutes the entire class. In other words: the set of functions computed by machines in $\hat{\mathcal{M}}$ is equal to the set of functions M^O , where $M \in \mathcal{M}$ and O is an arbitrary oracle. So when considering a code that is non-malleable with respect to a closed class \mathcal{F} relative to a random oracle model, it does not make a difference whether we give the tampering functions access to the random oracle or not. Summarizing the above discussion, we get the following theorem.

THEOREM 5.4. *Suppose $\mathcal{M}, \varepsilon, \rho, n, k$ such that \mathcal{M} is a set of closed oracle machines and $\log(\log(|\mathcal{M}|))$ is greater than the right-hand side of Equation (5) (see Theorem 5.1). Then there exists an efficient strongly non-malleable code with respect to \mathcal{M} in the random oracle model with exact security $\varepsilon + \rho$ even against an adversary making an unbounded number of queries to the random oracle.*

Examples of Closed Classes. An interesting and natural closed family consists of functions that tamper each half of the codeword independently: that is, $f \in \mathcal{F}_{\text{half}}$ if we can write $f : \{0, 1\}^n \rightarrow \{0, 1\}^n$ as $f(c_1, c_2) \stackrel{\text{def}}{=} (f_1(c_1), f_2(c_2))$ for $f_1, f_2 : \{0, 1\}^{n/2} \rightarrow \{0, 1\}^{n/2}$. This family has size $\log(\log(|\mathcal{F}_{\text{half}}|)) = n/2 + \log(n)$ as required by Theorem 5.1, and thus by Theorem 5.4, there exists an efficient strongly non-malleable code in the random oracle model with respect to $\mathcal{F}_{\text{half}}$ with rate $k/n \approx 1/6$.

More generally, for any $\delta < 1$, we can consider the class \mathcal{F}_δ where $f \in \mathcal{F}_\delta$ if every bit of $f(c)$ depends only on a subset of at most $\lfloor \delta n \rfloor$ bits of c . ($\mathcal{F}_{\text{half}}$ just discussed is thus a subset of $\mathcal{F}_{0.5}$.) The size of this class is $\log(\log(|\mathcal{F}_\delta|)) \leq \delta \cdot n + 2 \log(n)$ as required by Theorem 5.1.

6 TAMPER-RESILIENT SECURITY

In this section, we formally define the notion of tamper-resilient security as outlined in the Introduction. We consider some *interactive stateful system* $\langle G, s \rangle$ consisting of a *public* (possibly randomized) functionality $G : \{0, 1\}^u \times \{0, 1\}^n \rightarrow \{0, 1\}^v \times \{0, 1\}^n$ and *secret* initial state $s \in \{0, 1\}^n$. We consider two ways of interacting with the system:

Execute(x): A user can provide the system with **Execute(x)** queries, for $x \in \{0, 1\}^u$, in which case the system computes $(y, s') \leftarrow G(x, s)$, updates the state of the system to $s := s'$, and outputs y .

Tamper(f): We also consider tampering attacks against the system, modeled by **Tamper(f)** commands, for functions $f : \{0, 1\}^n \rightarrow \{0, 1\}^n$. Upon receiving such command, the system state is set to $s := f(s)$.

Honest users only interact with a system via **Execute** queries. As we discussed in the Introduction, an attacker that can also interact with the system via **Tamper** queries can potentially learn significantly more about the secret state, and even recover it entirely. Therefore, we would like to have a general method for securing systems against tampering attacks, so that the ability to issue **Tamper** queries (at least for functions f in some large family \mathcal{F}) cannot provide the attacker with additional information. We show how to use non-malleable codes for this purpose. First, we define what it means to *harden* a functionality G via a code (Enc, Dec) .

Definition 6.1. Let (Enc, Dec) be any coding scheme with k -bit messages and n -bit codewords. Let $G : \{0, 1\}^u \times \{0, 1\}^k \rightarrow \{0, 1\}^v \times \{0, 1\}^k$ be an arbitrary functionality with k -bit state. We define the *hardened functionality* $G^{\text{Enc}, \text{Dec}} : \{0, 1\}^u \times \{0, 1\}^n \rightarrow \{0, 1\}^v \times \{0, 1\}^n$ to be the functionality (with n -bit state) which, on input $(x, c) \in \{0, 1\}^u \times \{0, 1\}^n$, computes $s \leftarrow \text{Dec}(c)$ and checks if $s \stackrel{?}{=} \perp$. If so, it outputs a dummy value $(0^u, 0^n)$. Otherwise, it computes $(y, s') \leftarrow G(x, s)$ and outputs $(y, \text{Enc}(s'))$.

It is easy to see that the original system $\langle G, s \rangle$ and the hardened system $\langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle$ always behave exactly the same with respect to any series of **Execute** commands. Of course, they do *not* act the same with respect to **Tamper** commands, and that is the point; we wish to show that the hardened system $\langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle$ is tamper resilient and renders tampering attacks useless. We formally define this property next, as essentially saying that for any adversary \mathcal{A} that interacts with the hardened system $\langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle$ via **Execute** and **Tamper** queries, there is a simulator \mathcal{S} that interacts with the original system $\langle G, s \rangle$ *only via* **Execute** queries and learns the same information as \mathcal{A} . In particular, we say that a coding scheme (Enc, Dec) is *tamper simulatable* if it can be used to securely harden *any* system $\langle G, s \rangle$ into a tamper-resilient system $\langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle$.

Definition 6.2 (Tamper-Simulatability). A coding scheme (Enc, Dec) with k -bit messages and n -bit codewords is *tamper-simulatable with respect to a function family \mathcal{F}* , if there exists a simulator \mathcal{S} such that, for any functionality G with k bit state, any adversary \mathcal{A} , and any initial state $s \in \{0, 1\}^k$ we have

$$\text{TamperInteract}(\mathcal{A}, \mathcal{F}, \langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle) \approx \text{BBInteract}(\mathcal{S}, \langle G, s \rangle),$$

where **TamperInteract** and **BBInteract** are defined as follows.

TamperInteract $(\mathcal{A}, \mathcal{F}, \langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle)$: The adversary \mathcal{A} interacts with the system $\langle G, s \rangle$ for arbitrarily many rounds of interactions where, in each round:

- (1) The adversary can “tamper” by executing a **Tamper(f)** command against the system, for some $f \in \mathcal{F}$.
- (2) The adversary runs an **Execute(x)** query for some $x \in \{0, 1\}^u$, and receives the output y .

The output of the game consists of the output of the adversary \mathcal{A} at the end of the interaction, along with all of the **execute** queries x_1, x_2, \dots .

BBInteract $(\mathcal{S}, \langle G, s \rangle)$: The simulator interacts with the system $\langle G, s \rangle$ for arbitrarily many rounds of interaction where, in each round, it runs an **Execute(x)** query for some $x \in \{0, 1\}^u$ and receives the output y . The output of the game consists of the output of the simulator \mathcal{S} at the end of the interaction, along with all of the **execute**-query inputs x .

The simulator \mathcal{S} gets black-box access to the adversary \mathcal{A} , the functionality $G(\cdot, \cdot)$, and any tampering functions $f(\cdot)$ produced by \mathcal{A} (but does not get the state s). Here \approx can refer to statistical or computational indistinguishability.

Note that the `TamperInteract` and `BBInteract` games include all of the “execute” queries submitted by \mathcal{A} and \mathcal{S} in their outputs. This prevents the simulator from submitting queries which \mathcal{A} would not have asked: For example, if \mathcal{A} performs tampering attacks but only queries the system on some three inputs x_1, x_2, x_3 , then \mathcal{S} must simulate this without tampering queries *by only querying its system on the same inputs* x_1, x_2, x_3 . Also, note that in the definition of the `TamperInteract` game, we only allow the adversary to submit *one* tampering query `Tamper(f)` in each round. For families \mathcal{F} which are closed under composition (i.e., for any $f_1, f_2 \in \mathcal{F}$, we have $f_1 \circ f_2 \in \mathcal{F}$) this is without loss of generality, as submitting several tamper queries in a row can always be subsumed via a single query. We now show that a code (Enc, Dec) is tamper-simulatable with respect to \mathcal{F} if it is non-malleable with respect to \mathcal{F} .

THEOREM 6.1. *Let (Enc, Dec) be any coding scheme which is non-malleable with respect to \mathcal{F} . Then (Enc, Dec) is also tamper-simulatable with respect to \mathcal{F} .*

PROOF. By the definition of non-malleability, for each $f \in \mathcal{F}$ there exists a distribution D_f such that

$$\left\{ \begin{array}{l} c \leftarrow \text{Enc}(s), \tilde{c} \leftarrow f(c), \tilde{s} = \text{Dec}(\tilde{c}) \\ \text{Output: } \tilde{s}. \end{array} \right\} \approx \left\{ \begin{array}{l} \tilde{s} \leftarrow D_f \\ \text{Output } s \text{ if } \tilde{s} = \underline{\text{same}}, \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\}. \quad (15)$$

We use these distributions to define the simulator \mathcal{S} for the tamper simulatability definition. Recall that \mathcal{S} has black-box access to the adversary \mathcal{A} , the functionality $G(\cdot, \cdot)$, and the tampering function $f(\cdot)$ queried by \mathcal{A} . The simulation proceeds by running \mathcal{A} where, in each round:

- (1) When the adversary \mathcal{A} issues a `Tamper` command with function $f \in \mathcal{F}$ the simulator samples $\tilde{s} \leftarrow D_f$ (this can be done efficiently to give oracle access to $f(\cdot)$, by the definition to non-malleability).
- (2) When the adversary issues an `Execute` command with input x :
 - If $\tilde{s} = \underline{\text{same}}$, then the simulator \mathcal{S} forwards the `Execute(x)` query to its system $\langle G, s \rangle$ and forwards the output y to \mathcal{A} .
 - If $\tilde{s} \neq \underline{\text{same}}$, the simulator \mathcal{S} goes into “Overwritten Mode” (step 3).
- (3) “Overwritten Mode”: The simulator \mathcal{S} uses the modified state \tilde{s} to perfectly simulate the system $\langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(\tilde{s}) \rangle$ for all `Tamper` and `Execute` queries in all future rounds. Note that this is possible, using oracle access to $G(\cdot, \cdot)$ and to the tampering functions $f(\cdot)$.

To show the indistinguishability of simulation, we notice that, in each round *prior to and including* the round where the simulation enters “overwritten mode,” if the starting state in that round is s , the modification function is f and the input queried by the adversary is x , then (by the definition of non-malleability)

$$\begin{aligned} & \left\{ \begin{array}{l} c \leftarrow \text{Enc}(s), \tilde{c} \leftarrow f(c), \tilde{s} = \text{Dec}(\tilde{c}) \\ \text{Output: } G(x, \tilde{s}). \end{array} \right\} \\ & \approx \left\{ \begin{array}{l} \tilde{s} \leftarrow D_f \\ \text{Output: } G(x, s) \text{ if } \tilde{s} = \underline{\text{same}}, \text{ and } G(x, \tilde{s}) \text{ otherwise.} \end{array} \right\} \end{aligned}$$

where the distributions represent the output y seen by the adversary, and the updated “effective state” s of the system in the `TamperInteract` game (on the left-hand side) and in the simulation within the `BBInteract` game (on the right-hand side). Also, in each round *after* the simulation enters “overwritten mode” it acts exactly the same as the `TamperInteract` game. These two facts,

together with a hybrid argument over the number of rounds, prove the indistinguishability of the simulation.

More formally, let q be an upper bound on the number of tampering queries the adversary makes. We define hybrid distributions H_0, \dots, H_q where the i th hybrid H_i is defined by running BBInteract for the first i tamper queries, and then switch to TamperInteract , in particular, we will have

$$H_0 = \text{TamperInteract}(\mathcal{A}, \mathcal{F}, \langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle) \text{ and } H_q = \text{BBInteract}(\mathcal{S}, \langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle).$$

Note that H_q is equivalent to $\text{BBInteract}(\mathcal{S}, \langle G, s \rangle)$ as without tampering queries $\langle G, s \rangle$ and $\langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle$ are equivalent.

H_i is defined exactly as $\text{BBInteract}(\mathcal{S}, \langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(s) \rangle)$ until the adversary makes the i th tampering query $f(\cdot)$. If at this point the simulator is in overwritten mode, it just continues (as in step (3) above). Otherwise, the experiment switches to $\text{TamperInteract}(\mathcal{A}, \mathcal{F}, \langle G^{\text{Enc}, \text{Dec}}, \text{Enc}(\tilde{s}) \rangle)$ where $\text{Enc}(\tilde{s})$ is the current state. We claim that for any i , $H_i \approx H_{i+1}$, where \approx is as in Equation (15).¹¹ To see this, consider slightly different hybrids H'_i and H'_{i+1} , which are defined as H_i and H_{i+1} , but where after the i th tampering query we additionally output the state s' . Clearly, H'_i and H'_{i+1} are no harder to distinguish than H_i and H_{i+1} as one can always ignore the additional output. Moreover, we can w.l.o.g. assume the adversary stops after the i th tampering query as one can efficiently sample the rest of the experiment given s' . The outputs before the i th tampering query have the same distribution in H'_i and H'_{i+1} as up to this point the experiments are identical. So distinguishing H'_i from H'_{i+1} is equivalent to distinguishing the two distributions in Equation (15). \square

APPENDIX

A ALTERNATIVE DEFINITION OF NON-MALLEABLE CODES

In this section, we propose an alternative definition of non-malleability, inspired by an analogous definition for replayable CCA secure encryption in [15], which we show to be equivalent to our default notion. Essentially, the alternative definition says that we cannot distinguish the outcome of the tampering experiment that starts with source message s_0 from that which starts with source message s_1 , *unless* the outcome is one of $\{s_0, s_1\}$.

Definition A.1 (Alternative-Non-Malleability). Let \mathcal{F} be a family of functions. We say that an coding scheme (Enc, Dec) is *alternative-non-malleable with respect to \mathcal{F}* if for any $s_0, s_1 \in \{0, 1\}^k$ and any $f \in \mathcal{F}$, we have

$$\text{AltNM}_{s_0, s_1}^f(0) \approx \text{AltNM}_{s_0, s_1}^f(1),$$

where we define the two experiments by

$$\text{AltNM}_{s_0, s_1}^f(b) \stackrel{\text{def}}{=} \left\{ \begin{array}{l} c \leftarrow \text{Enc}(s_b), \tilde{c} \leftarrow f(c), \tilde{s} = \text{Dec}(\tilde{c}) \\ \text{Output } \underline{\text{same}} \text{ if } \tilde{s} \in \{s_0, s_1\}, \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\}$$

Here “ \approx ” can refer to statistical or computational indistinguishability. In the case of statistical indistinguishability, we say the scheme has error ε , if the statistical distance is at most ε .

We now show the equivalence of the two definitions of non-malleability. This equivalence only holds when the domain size of source messages is super-polynomial, as otherwise alternative-non-malleability may be a weaker notion. The following theorem makes this precise.

¹¹If we consider computational indistinguishability, we will lose a bit in the running time of the adversary considered. Concretely, if in Equation (15) we have security against adversaries running in time t , here we get security against time $t - t'$, where t' is the time required to invoke $G(\cdot, \cdot)$, $\text{Enc}(\cdot)$, $\text{Dec}(\cdot)$, and the simulator \mathcal{S} up to q times.

THEOREM A.1. For any function family \mathcal{F} :

- (1) If a coding scheme (Enc, Dec) is non-malleable with respect to \mathcal{F} , then it is alternatively non-malleable with respect to \mathcal{F} . In the case of statistical indistinguishability, if the scheme is non-malleable with error ϵ , then it is alternatively non-malleable with error 2ϵ .
- (2) If a coding scheme (Enc, Dec) is alternatively non-malleable with respect to \mathcal{F} , and its domain size is super-polynomial (i.e., $k = \omega(\log(\lambda))$ where λ is a security parameter), then it is also non-malleable with respect to \mathcal{F} . In the case of statistical indistinguishability, if the code is alternatively non-malleable with error ϵ , then it is non-malleable with error $\epsilon + 2^{-k}$.
- (3) When the domain size of the coding scheme is super-polynomial, a coding scheme (Enc, Dec) is non-malleable if and only if it is alternatively non-malleable.

PROOF. Let us start with the first part of the theorem. Assume (Enc, Dec) is non-malleable with respect to \mathcal{F} . Then, for each $f \in \mathcal{F}$ there exists some distribution D_f satisfying the definition of non-malleability. Therefore, for any $s_0, s_1 \in \{0, 1\}^k$,

$$\begin{aligned} \text{AltNM}_{s_0, s_1}^f(0) &\equiv \left\{ \begin{array}{l} c \leftarrow \text{Enc}(s_0), \tilde{c} \leftarrow f(c), \tilde{s} = \text{Dec}(\tilde{c}) \\ \text{Output } \underline{\text{same}} \text{ if } \tilde{s} \in \{s_0, s_1\}, \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\} \\ &\approx \left\{ \begin{array}{l} \tilde{s} \leftarrow D_f \\ \text{Output } \underline{\text{same}} \text{ if } \tilde{s} \in \{s_0, s_1, \underline{\text{same}}\} \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\} \\ &\approx \left\{ \begin{array}{l} c \leftarrow \text{Enc}(s_1), \tilde{c} \leftarrow f(c), \tilde{s} = \text{Dec}(\tilde{c}) \\ \text{Output } \underline{\text{same}} \text{ if } \tilde{s} \in \{s_0, s_1\}, \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\} \\ &\equiv \text{AltNM}_{s_0, s_1}^f(0). \end{aligned}$$

In the case of statistical indistinguishability, each “ \approx ” above hides a factor of ϵ statistical distance, giving us a total of 2ϵ .

For the second part of the theorem, assume that (Enc, Dec) is alternatively non-malleable with respect to \mathcal{F} . We construct the distribution D_f as follows:

$$D_f \stackrel{\text{def}}{=} \left\{ \begin{array}{l} \text{Sample } s \leftarrow \{0, 1\}^k, c \leftarrow \text{Enc}(s), \tilde{c} \leftarrow f(c), \tilde{s} = \text{Dec}(\tilde{c}) \\ \text{Output } \underline{\text{same}} \text{ if } \tilde{s} = s \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\}$$

In other words, D_f tries running the tampering experiment on a random s and outputs same if the value remains unchanged. We show that D_f satisfies the definition of non-malleability. For any $f \in \mathcal{F}$, $s_0 \in \{0, 1\}^k$,

$$\begin{aligned} &\left\{ \begin{array}{l} \tilde{s} \leftarrow D_f \\ \text{Output } s_0 \text{ if } \tilde{s} = \underline{\text{same}} \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\} \\ &\equiv \left\{ \begin{array}{l} s \leftarrow \{0, 1\}^k, \tilde{s} \leftarrow \text{AltNM}_{s_0, s}^f(1) \\ \text{Output } s_0 \text{ if } \tilde{s} = \underline{\text{same}} \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\} \\ &\approx \left\{ \begin{array}{l} s \leftarrow \{0, 1\}^k, \tilde{s} \leftarrow \text{AltNM}_{s_0, s}^f(0) \\ \text{Output } s_0 \text{ if } \tilde{s} = \underline{\text{same}} \text{ and } \tilde{s} \text{ otherwise.} \end{array} \right\} \\ &\approx \left\{ \begin{array}{l} c \leftarrow \text{Enc}(s_0), \tilde{c} \leftarrow f(c), \tilde{s} = \text{Dec}(\tilde{c}) \\ \text{Output } \tilde{s}. \end{array} \right\} \end{aligned}$$

For statistical indistinguishability, the first “ \approx ” hides a factor of ϵ while the second “ \approx ” (which is always statistical) hides a factor of 2^{-k} since the third and fourth experiments only differ when $\text{Dec}(\tilde{c}) = s$, where s is chosen randomly and independently of how \tilde{c} is generated, which only happens with probability 2^{-k} .

The third part of the theorem simply follows from the previous two. \square

ACKNOWLEDGMENTS

We would like to thank Ronald Cramer, Yevgeniy Dodis, Eike Kiltz, Swastik Kopparty, Ignacio Cascudo Pueyo and Madhu Sudan for the many helpful discussions and comments they provided for this project. We also thank the anonymous *Journal of the ACM* reviewers for their comments.

REFERENCES

- [1] Divesh Aggarwal. 2015. Affine-evasive sets modulo a prime. *Inf. Process. Lett.* 115, 2 (2015), 382–385. DOI: <http://dx.doi.org/10.1016/j.ipl.2014.10.015>
- [2] Divesh Aggarwal, Yevgeniy Dodis, Tomasz Kazana, and Maciej Obremski. 2015a. Non-malleable reductions and applications. In *47th Annual ACM Symposium on Theory of Computing*, Rocco A. Servedio and Ronitt Rubinfeld (Eds.). ACM, 459–468.
- [3] Divesh Aggarwal, Yevgeniy Dodis, and Shachar Lovett. 2014. Non-malleable codes from additive combinatorics. In *46th Annual ACM Symposium on Theory of Computing*, David B. Shmoys (Ed.). ACM, New York, 774–783.
- [4] Divesh Aggarwal, Stefan Dziembowski, Tomasz Kazana, and Maciej Obremski. 2015b. Leakage-resilient non-malleable codes. In *TCC 2015: 12th Theory of Cryptography Conference, Part I*, Yevgeniy Dodis and Jesper Buus Nielsen (Eds.). Lecture Notes in Computer Science, Vol. 9014, Springer, Heidelberg, 398–426. DOI: http://dx.doi.org/10.1007/978-3-662-46494-6_17
- [5] Shashank Agrawal, Divya Gupta, Hemanta K. Maji, Omkant Pandey, and Manoj Prabhakaran. 2015a. Explicit non-malleable codes against bit-wise tampering and permutations. In *Advances in Cryptology – CRYPTO 2015, Part I*, Rosario Gennaro and Matthew J. B. Robshaw (Eds.). Lecture Notes in Computer Science, Vol. 9215, Springer, Heidelberg, 538–557. DOI: http://dx.doi.org/10.1007/978-3-662-47989-6_26
- [6] Shashank Agrawal, Divya Gupta, Hemanta K. Maji, Omkant Pandey, and Manoj Prabhakaran. 2015b. A rate-optimizing compiler for non-malleable codes against bit-wise tampering and permutations. In *TCC 2015: 12th Theory of Cryptography Conference, Part I*, Yevgeniy Dodis and Jesper Buus Nielsen (Eds.). Lecture Notes in Computer Science, Vol. 9014, Springer, Heidelberg, 375–397. DOI: http://dx.doi.org/10.1007/978-3-662-46494-6_16
- [7] Adi Akavia, Shafi Goldwasser, and Vinod Vaikuntanathan. 2009. Simultaneous hardcore bits and cryptography against memory attacks. In *TCC 2009: 6th Theory of Cryptography Conference*, Omer Reingold (Ed.). Lecture Notes in Computer Science, Vol. 5444, Springer, Heidelberg, 474–495.
- [8] Joël Alwen, Yevgeniy Dodis, and Daniel Wichs. 2009. Leakage-resilient public-key cryptography in the bounded-retrieval model. In *Advances in Cryptology – CRYPTO 2009*, Shai Halevi (Ed.). Lecture Notes in Computer Science, Vol. 5677, Springer, Heidelberg, 36–54.
- [9] Mihir Bellare and Chanathip Namprempre. 2000. Authenticated encryption: Relations among notions and analysis of the generic composition paradigm. In *Advances in Cryptology – ASIACRYPT 2000*, Tatsuaki Okamoto (Ed.). Lecture Notes in Computer Science, Vol. 1766, Springer, Heidelberg, 531–545.
- [10] Mihir Bellare and Phillip Rogaway. 1993. Random oracles are practical: A paradigm for designing efficient protocols. In *ACM CCS 93: 1st Conference on Computer and Communications Security*, V. Ashby (Ed.). ACM, 62–73.
- [11] Mihir Bellare and John Rompel. 1994. Randomness-efficient oblivious sampling. In *35th Annual Symposium on Foundations of Computer Science*. IEEE Computer Society Press, 276–287.
- [12] Dan Boneh, Richard A. DeMillo, and Richard J. Lipton. 1997. On the importance of checking cryptographic protocols for faults (extended abstract). In *Advances in Cryptology – EUROCRYPT’97*, Walter Fumy (Ed.). Lecture Notes in Computer Science, Vol. 1233, Springer, Heidelberg, 37–51.
- [13] Zvika Brakerski, Yael Tauman Kalai, Jonathan Katz, and Vinod Vaikuntanathan. 2010. Overcoming the hole in the bucket: Public-key cryptography resilient to continual memory leakage. In *51st Annual Symposium on Foundations of Computer Science*. IEEE Computer Society Press, 501–510.
- [14] Sergio Cabello, Carles Padró, and Germán Sáez. 2002. Secret sharing schemes with detection of cheaters for a general access structure. *Des. Codes Cryptography* 25, 2 (2002), 175–188.
- [15] Ran Canetti, Hugo Krawczyk, and Jesper Buus Nielsen. 2003. Relaxing chosen-ciphertext security. In *Advances in Cryptology – CRYPTO 2003*, Dan Boneh (Ed.). Lecture Notes in Computer Science, Vol. 2729, Springer, Heidelberg, 565–582.
- [16] Hervé Chabanne, Gérard Cohen, Jean-Pierre Flori, and Alain Patey. 2011. Non-malleable codes from the wire-tap channel. In *2011 IEEE Information Theory Workshop (ITW’11)*. 55–59. DOI: <http://dx.doi.org/10.1109/ITW.2011.6089565>
- [17] Hervé Chabanne, Gérard D. Cohen, and Alain Patey. 2012. Secure network coding and non-malleable codes: Protection against linear tampering. In *2012 IEEE International Symposium on Information Theory (ISIT’12)*. IEEE, 2546–2550. DOI: <http://dx.doi.org/10.1109/ISIT.2012.6283976>
- [18] Nishanth Chandran, Vipul Goyal, Pratyay Mukherjee, Omkant Pandey, and Jalaj Upadhyay. 2016. Block-wise non-malleable codes. In *43rd International Colloquium on Automata, Languages, and Programming (ICALP’16) (LIPIcs)*,

- Ioannis Chatzigiannakis, Michael Mitzenmacher, Yuval Rabani, and Davide Sangiorgi (Eds.), Vol. 55. Schloss Dagstuhl - Leibniz-Zentrum fuer Informatik, 31:1–31:14. DOI : <http://dx.doi.org/10.4230/LIPIcs.ICALP.2016.31>
- [19] Suresh Chari, Charanjit S. Jutla, Josyula R. Rao, and Pankaj Rohatgi. 1999. Towards sound approaches to counteract power-analysis attacks. In *Advances in Cryptology – CRYPTO’99*, Michael J. Wiener (Ed.). Lecture Notes in Computer Science, Vol. 1666, Springer, Heidelberg, 398–412.
 - [20] Eshan Chattopadhyay and David Zuckerman. 2014. Non-malleable codes against constant split-state tampering. In *55th IEEE Annual Symposium on Foundations of Computer Science (FOCS’14)*. IEEE Computer Society, 306–315. DOI : <http://dx.doi.org/10.1109/FOCS.2014.40>
 - [21] Hao Chen, Ronald Cramer, Shafi Goldwasser, Robbert de Haan, and Vinod Vaikuntanathan. 2007. Secure computation from random error correcting codes. In *Advances in Cryptology – EUROCRYPT 2007*, Moni Naor (Ed.). Lecture Notes in Computer Science, Vol. 4515, Springer, Heidelberg, 291–310.
 - [22] Mahdi Cheraghchi and Venkatesan Guruswami. 2014a. Capacity of non-malleable codes. In *Innovations in Theoretical Computer Science (ITCS’14)*, Moni Naor (Ed.). ACM, 155–168. DOI : <http://dx.doi.org/10.1145/2554797.2554814>
 - [23] Mahdi Cheraghchi and Venkatesan Guruswami. 2014b. Non-malleable coding against bit-wise and split-state tampering. In *TCC 2014: 11th Theory of Cryptography Conference*, Yehuda Lindell (Ed.). Lecture Notes in Computer Science, Vol. 8349, Springer, Heidelberg, 440–464. DOI : http://dx.doi.org/10.1007/978-3-642-54242-8_19
 - [24] Sandro Coretti, Yevgeniy Dodis, Björn Tackmann, and Daniele Venturi. 2016. Non-malleable encryption: Simpler, shorter, stronger. In *TCC 2016-A: 13th Theory of Cryptography Conference, Part I*, Eyal Kushilevitz and Tal Malkin (Eds.). Lecture Notes in Computer Science, Vol. 9562, Springer, Heidelberg, 306–335. DOI : http://dx.doi.org/10.1007/978-3-662-49096-9_13
 - [25] Sandro Coretti, Ueli Maurer, Björn Tackmann, and Daniele Venturi. 2015. From single-bit to multi-bit public-key encryption via non-malleable codes. In *TCC 2015: 12th Theory of Cryptography Conference, Part I*, Yevgeniy Dodis and Jesper Buus Nielsen (Eds.). Lecture Notes in Computer Science, Vol. 9014, Springer, Heidelberg, 532–560. DOI : http://dx.doi.org/10.1007/978-3-662-46494-6_22
 - [26] Jean-Sébastien Coron and Louis Goubin. 2000. On boolean and arithmetic masking against differential power analysis. In *Cryptographic Hardware and Embedded Systems – CHES 2000*, Çetin Kaya Koç and Christof Paar (Eds.). Lecture Notes in Computer Science, Vol. 1965, Springer, Heidelberg, 231–237.
 - [27] Jean-Sébastien Coron, Thomas Holenstein, Robin Künzler, Jacques Patarin, Yannick Seurin, and Stefano Tessaro. 2016. How to build an ideal cipher: The indistinguishability of the Feistel construction. *Journal of Cryptology* 29, 1 (Jan. 2016), 61–114. DOI : <http://dx.doi.org/10.1007/s00145-014-9189-6>
 - [28] Ronald Cramer, Yevgeniy Dodis, Serge Fehr, Carles Padró, and Daniel Wichs. 2008. Detection of algebraic manipulation with applications to robust secret sharing and fuzzy extractors. In *Advances in Cryptology – EUROCRYPT 2008*, Nigel P. Smart (Ed.). Lecture Notes in Computer Science, Vol. 4965, Springer, Heidelberg, 471–488.
 - [29] Dana Dachman-Soled, Feng-Hao Liu, Elaine Shi, and Hong-Sheng Zhou. 2015. Locally decodable and updatable non-malleable codes and their applications. In *TCC 2015: 12th Theory of Cryptography Conference, Part I*, Yevgeniy Dodis and Jesper Buus Nielsen (Eds.). Lecture Notes in Computer Science, Vol. 9014, Springer, Heidelberg, 427–450. DOI : http://dx.doi.org/10.1007/978-3-662-46494-6_18
 - [30] Francesco Davi, Stefan Dziembowski, and Daniele Venturi. 2010. Leakage-resilient storage. In *SCN 10: 7th International Conference on Security in Communication Networks*, Juan A. Garay and Roberto De Prisco (Eds.). Lecture Notes in Computer Science, Vol. 6280, Springer, Heidelberg, 121–137.
 - [31] Yevgeniy Dodis, Kristiyan Haralambiev, Adriana López-Alt, and Daniel Wichs. 2010. Cryptography against continuous memory attacks. In *51st Annual Symposium on Foundations of Computer Science*. IEEE Computer Society Press, 511–520.
 - [32] Yevgeniy Dodis, Yael Tauman Kalai, and Shachar Lovett. 2009. On cryptography with auxiliary input. In *41st Annual ACM Symposium on Theory of Computing*, Michael Mitzenmacher (Ed.). ACM, 621–630.
 - [33] Yevgeniy Dodis, Allison B. Lewko, Brent Waters, and Daniel Wichs. 2011. Storing secrets on continually leaky devices. In *52nd Annual Symposium on Foundations of Computer Science*, Rafail Ostrovsky (Ed.). IEEE Computer Society Press, 688–697.
 - [34] Danny Dolev, Cynthia Dwork, and Moni Naor. 2000. Nonmalleable cryptography. *SIAM J. Comput.* 30, 2 (2000), 391–437.
 - [35] Alexandre Duc, Stefan Dziembowski, and Sebastian Faust. 2014. Unifying leakage models: From probing attacks to noisy leakage. In *Advances in Cryptology – EUROCRYPT 2014*, Phong Q. Nguyen and Elisabeth Oswald (Eds.). Lecture Notes in Computer Science, Vol. 8441, Springer, Heidelberg, 423–440. DOI : http://dx.doi.org/10.1007/978-3-642-55220-5_24
 - [36] Stefan Dziembowski and Sebastian Faust. 2011. Leakage-resilient cryptography from the inner-product extractor. In *Advances in Cryptology – ASIACRYPT 2011*, Dong Hoon Lee and Xiaoyun Wang (Eds.). Lecture Notes in Computer Science, Vol. 7073, Springer, Heidelberg, 702–721.

- [37] Stefan Dziembowski and Sebastian Faust. 2012. Leakage-resilient circuits without computational assumptions. In *TCC 2012: 9th Theory of Cryptography Conference*, Ronald Cramer (Ed.). Lecture Notes in Computer Science, Vol. 7194, Springer, Heidelberg, 230–247.
- [38] Stefan Dziembowski, Tomasz Kazana, and Maciej Obremski. 2013. Non-malleable codes from two-source extractors. In *Advances in Cryptology – CRYPTO 2013, Part II*, Ran Canetti and Juan A. Garay (Eds.). Lecture Notes in Computer Science, Vol. 8043, Springer, Heidelberg, 239–257. DOI : http://dx.doi.org/10.1007/978-3-642-40084-1_14
- [39] Stefan Dziembowski and Krzysztof Pietrzak. 2008. Leakage-resilient cryptography. In *49th Annual Symposium on Foundations of Computer Science*. IEEE Computer Society Press, 293–302.
- [40] Stefan Dziembowski, Krzysztof Pietrzak, and Daniel Wichs. 2010. Non-malleable codes. In *ICS 2010: 1st Innovations in Computer Science*, Andrew Chi-Chih Yao (Ed.). Tsinghua University Press, Tsinghua University, Beijing, China, 434–452.
- [41] Sebastian Faust, Eike Kiltz, Krzysztof Pietrzak, and Guy N. Rothblum. 2010. Leakage-resilient signatures. In *TCC 2010: 7th Theory of Cryptography Conference*, Daniele Micciancio (Ed.). Lecture Notes in Computer Science, Vol. 5978, Springer, Heidelberg, 343–360.
- [42] Sebastian Faust, Pratyay Mukherjee, Jesper Buus Nielsen, and Daniele Venturi. 2014a. Continuous non-malleable codes. In *TCC 2014: 11th Theory of Cryptography Conference*, Yehuda Lindell (Ed.). Lecture Notes in Computer Science, Vol. 8349, Springer, Heidelberg, 465–488. DOI : http://dx.doi.org/10.1007/978-3-642-54242-8_20
- [43] Sebastian Faust, Pratyay Mukherjee, Daniele Venturi, and Daniel Wichs. 2014b. Efficient non-malleable codes and key-derivation for poly-size tampering circuits. In *Advances in Cryptology – EUROCRYPT 2014*, Phong Q. Nguyen and Elisabeth Oswald (Eds.). Lecture Notes in Computer Science, Vol. 8441, Springer, Heidelberg, 111–128. DOI : http://dx.doi.org/10.1007/978-3-642-55220-5_7
- [44] Sebastian Faust, Tal Rabin, Leonid Reyzin, Eran Tromer, and Vinod Vaikuntanathan. 2010. Protecting circuits from leakage: The computationally-bounded and noisy cases. In *Advances in Cryptology – EUROCRYPT 2010*, Henri Gilbert (Ed.). Lecture Notes in Computer Science, Vol. 6110, Springer, Heidelberg, Germany, 135–156.
- [45] Rosario Gennaro, Anna Lysyanskaya, Tal Malkin, Silvio Micali, and Tal Rabin. 2004. Algorithmic tamper-proof (ATP) security: Theoretical foundations for security against hardware tampering. In *TCC 2004: 1st Theory of Cryptography Conference*, Moni Naor (Ed.). Lecture Notes in Computer Science, Vol. 2951, Springer, Heidelberg, 258–277.
- [46] Louis Goubin and Jacques Patarin. 1999. DES and differential power analysis (The “duplication” method). In *Cryptographic Hardware and Embedded Systems – CHES’99*, Çetin Kaya Koç and Christof Paar (Eds.). Lecture Notes in Computer Science, Vol. 1717, Springer, Heidelberg, 158–172.
- [47] Shai Halevi and Huijia Lin. 2011. After-the-fact leakage in public-key encryption. In *TCC 2011: 8th Theory of Cryptography Conference*, Yuval Ishai (Ed.). Lecture Notes in Computer Science, Vol. 6597, Springer, Heidelberg, 107–124.
- [48] Richard W. Hamming. 1950. Error detecting and error correcting codes. *Bell System Technical Journal* 29 (1950), 147–160.
- [49] Carmit Hazay, Adriana López-Alt, Hoeteck Wee, and Daniel Wichs. 2013. Leakage-resilient cryptography from minimal assumptions. In *Advances in Cryptology – EUROCRYPT 2013*, Thomas Johansson and Phong Q. Nguyen (Eds.). Lecture Notes in Computer Science, Vol. 7881, Springer, Heidelberg, 160–176. DOI : http://dx.doi.org/10.1007/978-3-642-38348-9_10
- [50] Yuval Ishai, Manoj Prabhakaran, Amit Sahai, and David Wagner. 2006. Private circuits II: Keeping secrets in tamperable circuits. In *Advances in Cryptology – EUROCRYPT 2006*, Serge Vaudenay (Ed.). Lecture Notes in Computer Science, Vol. 4004, Springer, Heidelberg, 308–327.
- [51] Yuval Ishai, Amit Sahai, and David Wagner. 2003. Private circuits: Securing hardware against probing attacks. In *Advances in Cryptology – CRYPTO 2003*, Dan Boneh (Ed.). Lecture Notes in Computer Science, Vol. 2729, Springer, Heidelberg, 463–481.
- [52] Zahra Jafargholi and Daniel Wichs. 2015. Tamper detection and continuous non-malleable codes. In *TCC 2015: 12th Theory of Cryptography Conference, Part I*, Yevgeniy Dodis and Jesper Buus Nielsen (Eds.). Lecture Notes in Computer Science, Vol. 9014, Springer, Heidelberg, 451–480. DOI : http://dx.doi.org/10.1007/978-3-662-46494-6_19
- [53] Jonathan Katz and Vinod Vaikuntanathan. 2009. Signature schemes with bounded leakage resilience. In *Advances in Cryptology – ASIACRYPT 2009*, Mitsuru Matsui (Ed.). Lecture Notes in Computer Science, Vol. 5912, Springer, Heidelberg, 703–720.
- [54] Feng-Hao Liu and Anna Lysyanskaya. 2012. Tamper and leakage resilience in the split-state model. In *Advances in Cryptology – CRYPTO 2012*, Reihaneh Safavi-Naini and Ran Canetti (Eds.). Lecture Notes in Computer Science, Vol. 7417, Springer, Heidelberg, 517–532.
- [55] Ueli M. Maurer, Renato Renner, and Clemens Holenstein. 2004. Indifferentiability, impossibility results on reductions, and applications to the random oracle methodology. In *TCC 2004: 1st Theory of Cryptography Conference*, Moni Naor (Ed.). Lecture Notes in Computer Science, Vol. 2951, Springer, Heidelberg, 21–39.

- [56] Thomas S. Messerges. 2001. Securing the AES finalists against power analysis attacks. In *Fast Software Encryption – FSE 2000*, Bruce Schneier (Ed.). Lecture Notes in Computer Science, Vol. 1978, Springer, Heidelberg, 150–164.
- [57] Silvio Micali and Leonid Reyzin. 2004. Physically observable cryptography (extended abstract). In *TCC 2004: 1st Theory of Cryptography Conference*, Moni Naor (Ed.). Lecture Notes in Computer Science, Vol. 2951, Springer, Heidelberg, 278–296.
- [58] Michael Mitzenmacher. 2008. A survey of results for deletion channels and related synchronization channels. In *Algorithm Theory - SWAT 2008*, Joachim Gudmundsson (Ed.). Lecture Notes in Computer Science, Vol. 5124, Springer, Berlin, 1–3. http://dx.doi.org/10.1007/978-3-540-69903-3_1
- [59] Moni Naor and Gil Segev. 2009. Public-key cryptosystems resilient to key leakage. In *Advances in Cryptology – CRYPTO 2009*, Shai Halevi (Ed.). Lecture Notes in Computer Science, Vol. 5677 Springer, Heidelberg, 18–35.
- [60] Wakaha Ogata and Kaoru Kurosawa. 1996. Optimum secret sharing scheme secure against cheating. In *Advances in Cryptology – EUROCRYPT ’96*, Ueli M. Maurer (Ed.). Lecture Notes in Computer Science, Vol. 1070, Springer, Heidelberg, Germany, Saragossa, Spain, 200–211.
- [61] Wakaha Ogata, Kaoru Kurosawa, Douglas R. Stinson, and Hajime Saido. 2004. New combinatorial designs and their applications to authentication codes and secret sharing schemes. *Discrete Mathematics* 279, 1–3 (2004), 383–405.
- [62] Lawrence H. Ozarow and Aaron D. Wyner. 1985. Wire-tap channel II. In *Advances in Cryptology (EUROCRYPT’84)*, Thomas Beth, Norbert Cot, and Ingemar Ingemarsson (Eds.). Lecture Notes in Computer Science, Vol. 209, Springer, Heidelberg, 33–50.
- [63] Krzysztof Pietrzak. 2009. A leakage-resilient mode of operation. In *Advances in Cryptology (EUROCRYPT’09)*, Antoine Joux (Ed.). Lecture Notes in Computer Science, Vol. 5479, Springer, Heidelberg, 462–482.
- [64] Ignacio Cascudo Pueyo, Hao Chen, Ronald Cramer, and Chaoping Xing. 2009. Asymptotically good ideal linear secret sharing with strong multiplication over any fixed finite field. In *Advances in Cryptology – CRYPTO 2009*, Shai Halevi (Ed.). Lecture Notes in Computer Science, Vol. 5677, Springer, Heidelberg, 466–486.
- [65] Jeanette P. Schmidt, Alan Siegel, and Aravind Srinivasan. 1993. Chernoff-Hoeffding bounds for applications with limited independence. In *The 4th Annual ACM-SIAM Symposium on Discrete Algorithms (SODA ’93)*. Society for Industrial and Applied Mathematics, Philadelphia, 331–340.
- [66] Adi Shamir. 1979. How to share a secret. *Commun. Assoc. Comput. Machin.* 22, 11 (Nov. 1979), 612–613.
- [67] Claude E. Shannon. 1949. Communication theory of secrecy systems. *Bell Syst. Tech. J.* 28, 4 (1949), 656–715.

Received March 2015; revised December 2016; accepted January 2018